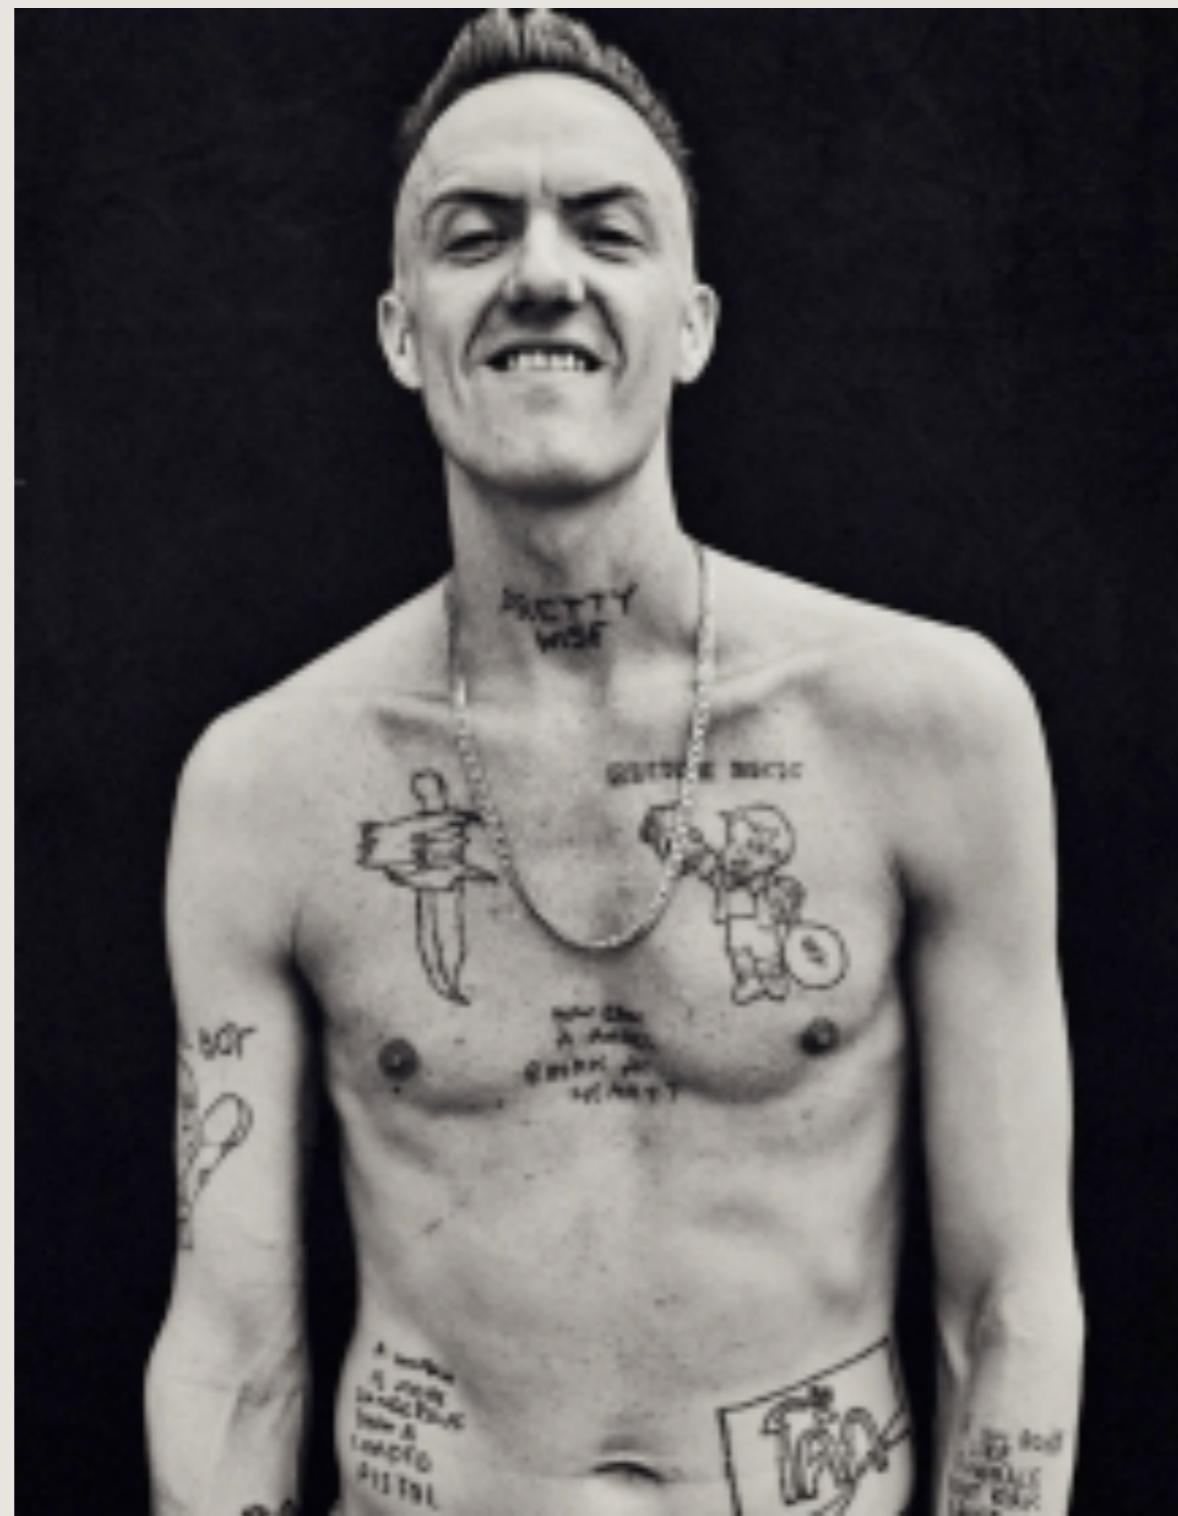


production monitoring. zef side.





stylefruits.de

here not to teach

but to bring your
attention

how many
servers do you
have?

how many page
views do you
have?

do you monitor

your system

extensively?

what do you use?

So,
production? ...

looks like that.

**yeah, that makes
total sense!...**

(c) noone. never.

what's there? ...

looking for
patterns

identify

constructs

repetitions

trends

outliers

dispersion

deviation

fragmentation

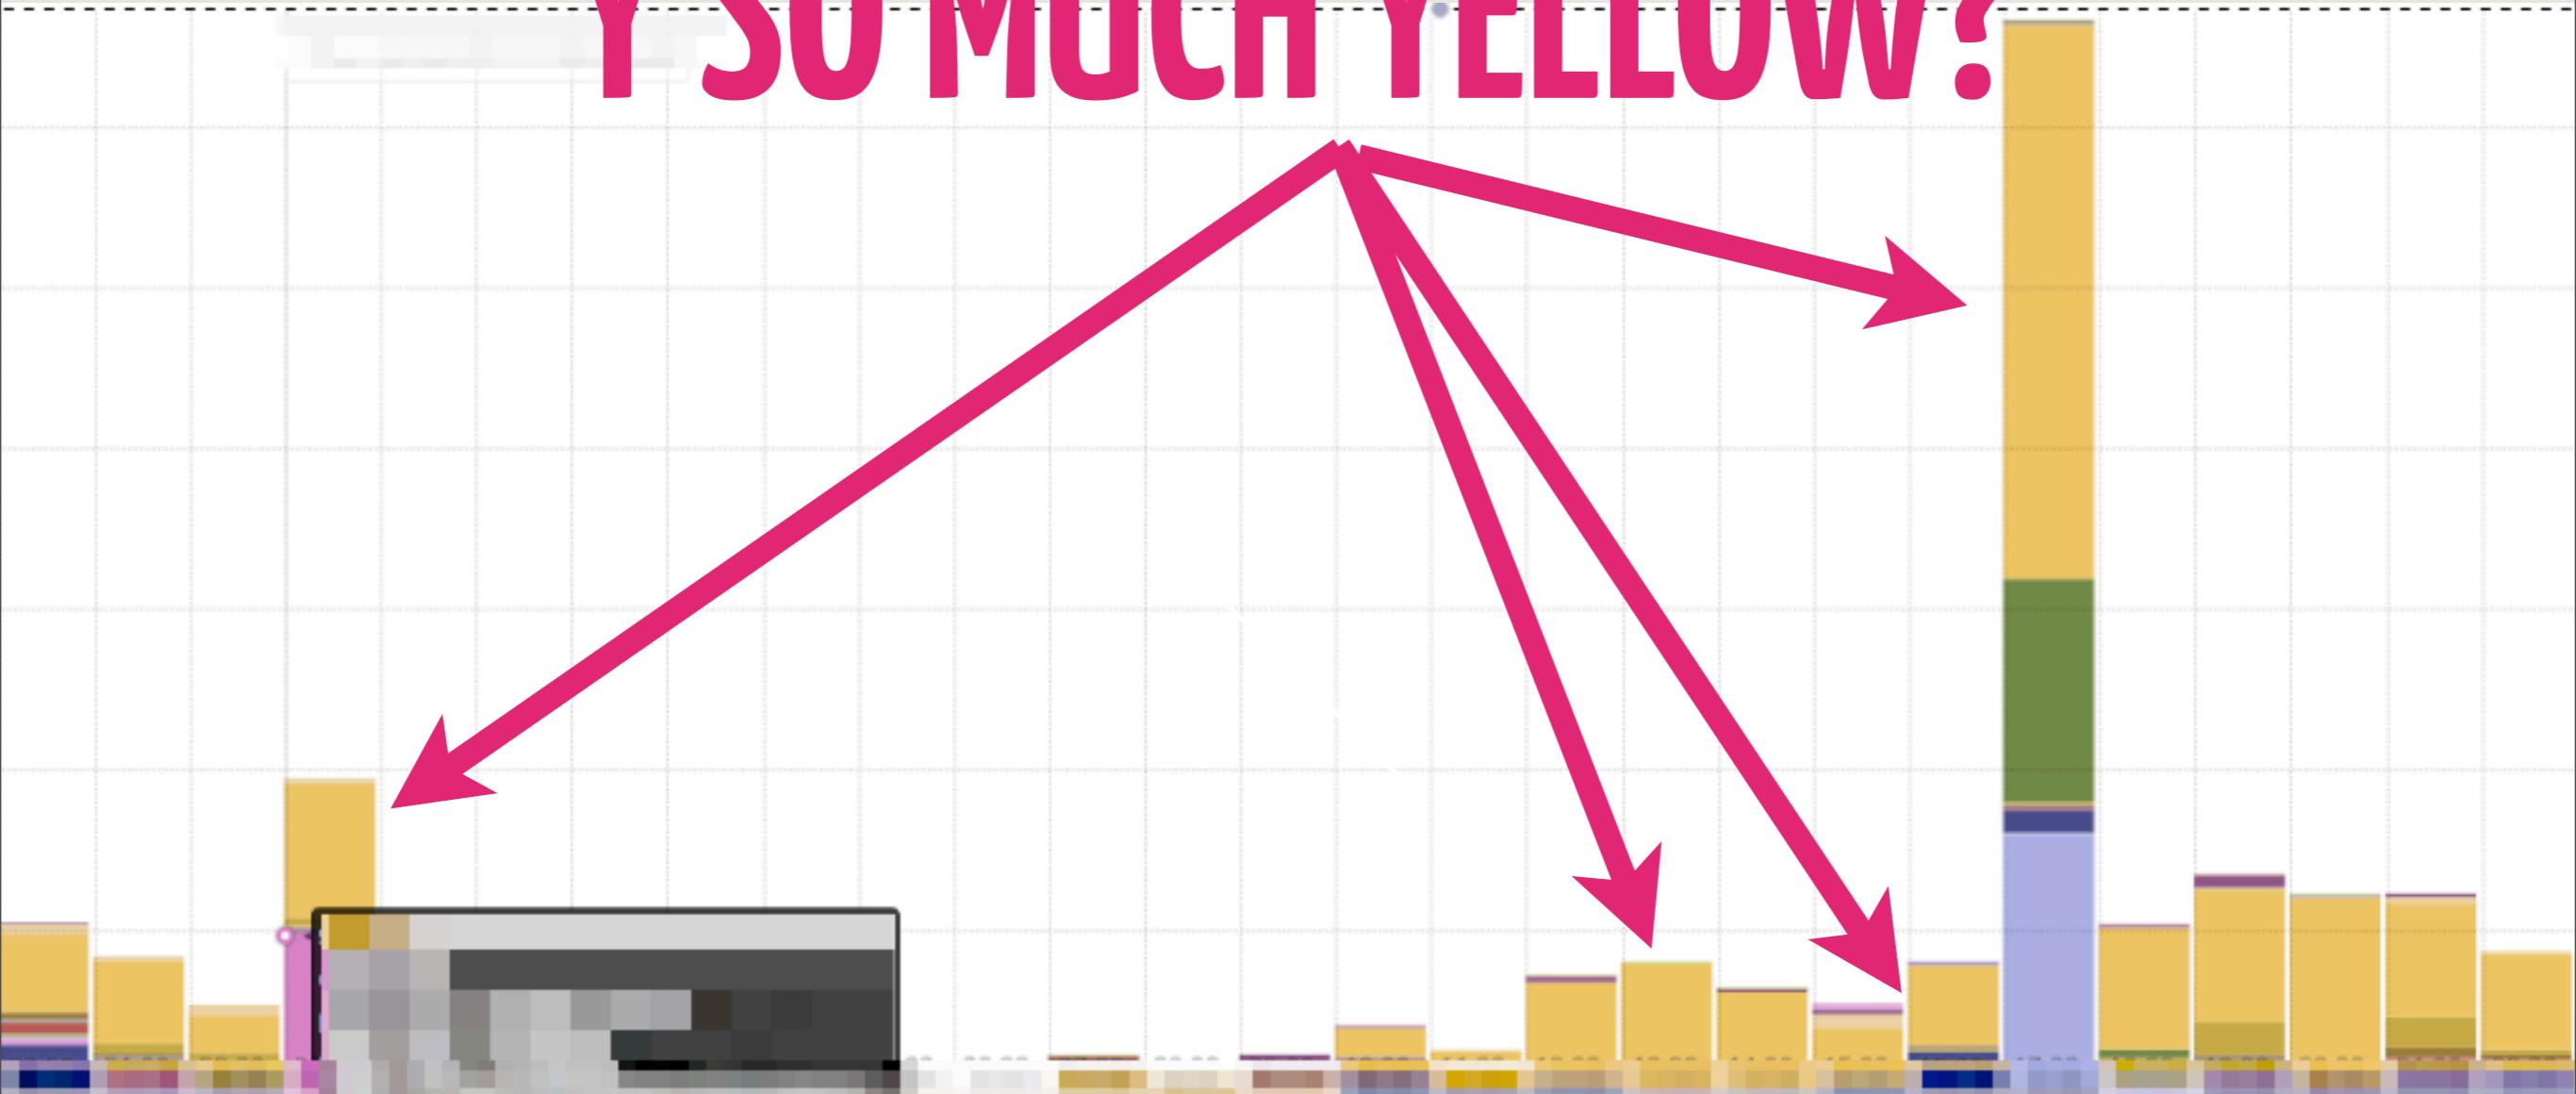
true story:

sittin' there,

gathering
exceptions from
production

suddenly

Y SO MUCH YELLOW?



check into
exception
details...

3 words:

parse-query-string

Ah shoo!

That's not really
an exception!

That's just a 404!

Eliminate the cause

**Avoid hitting complete stack when
there's no way to do it**

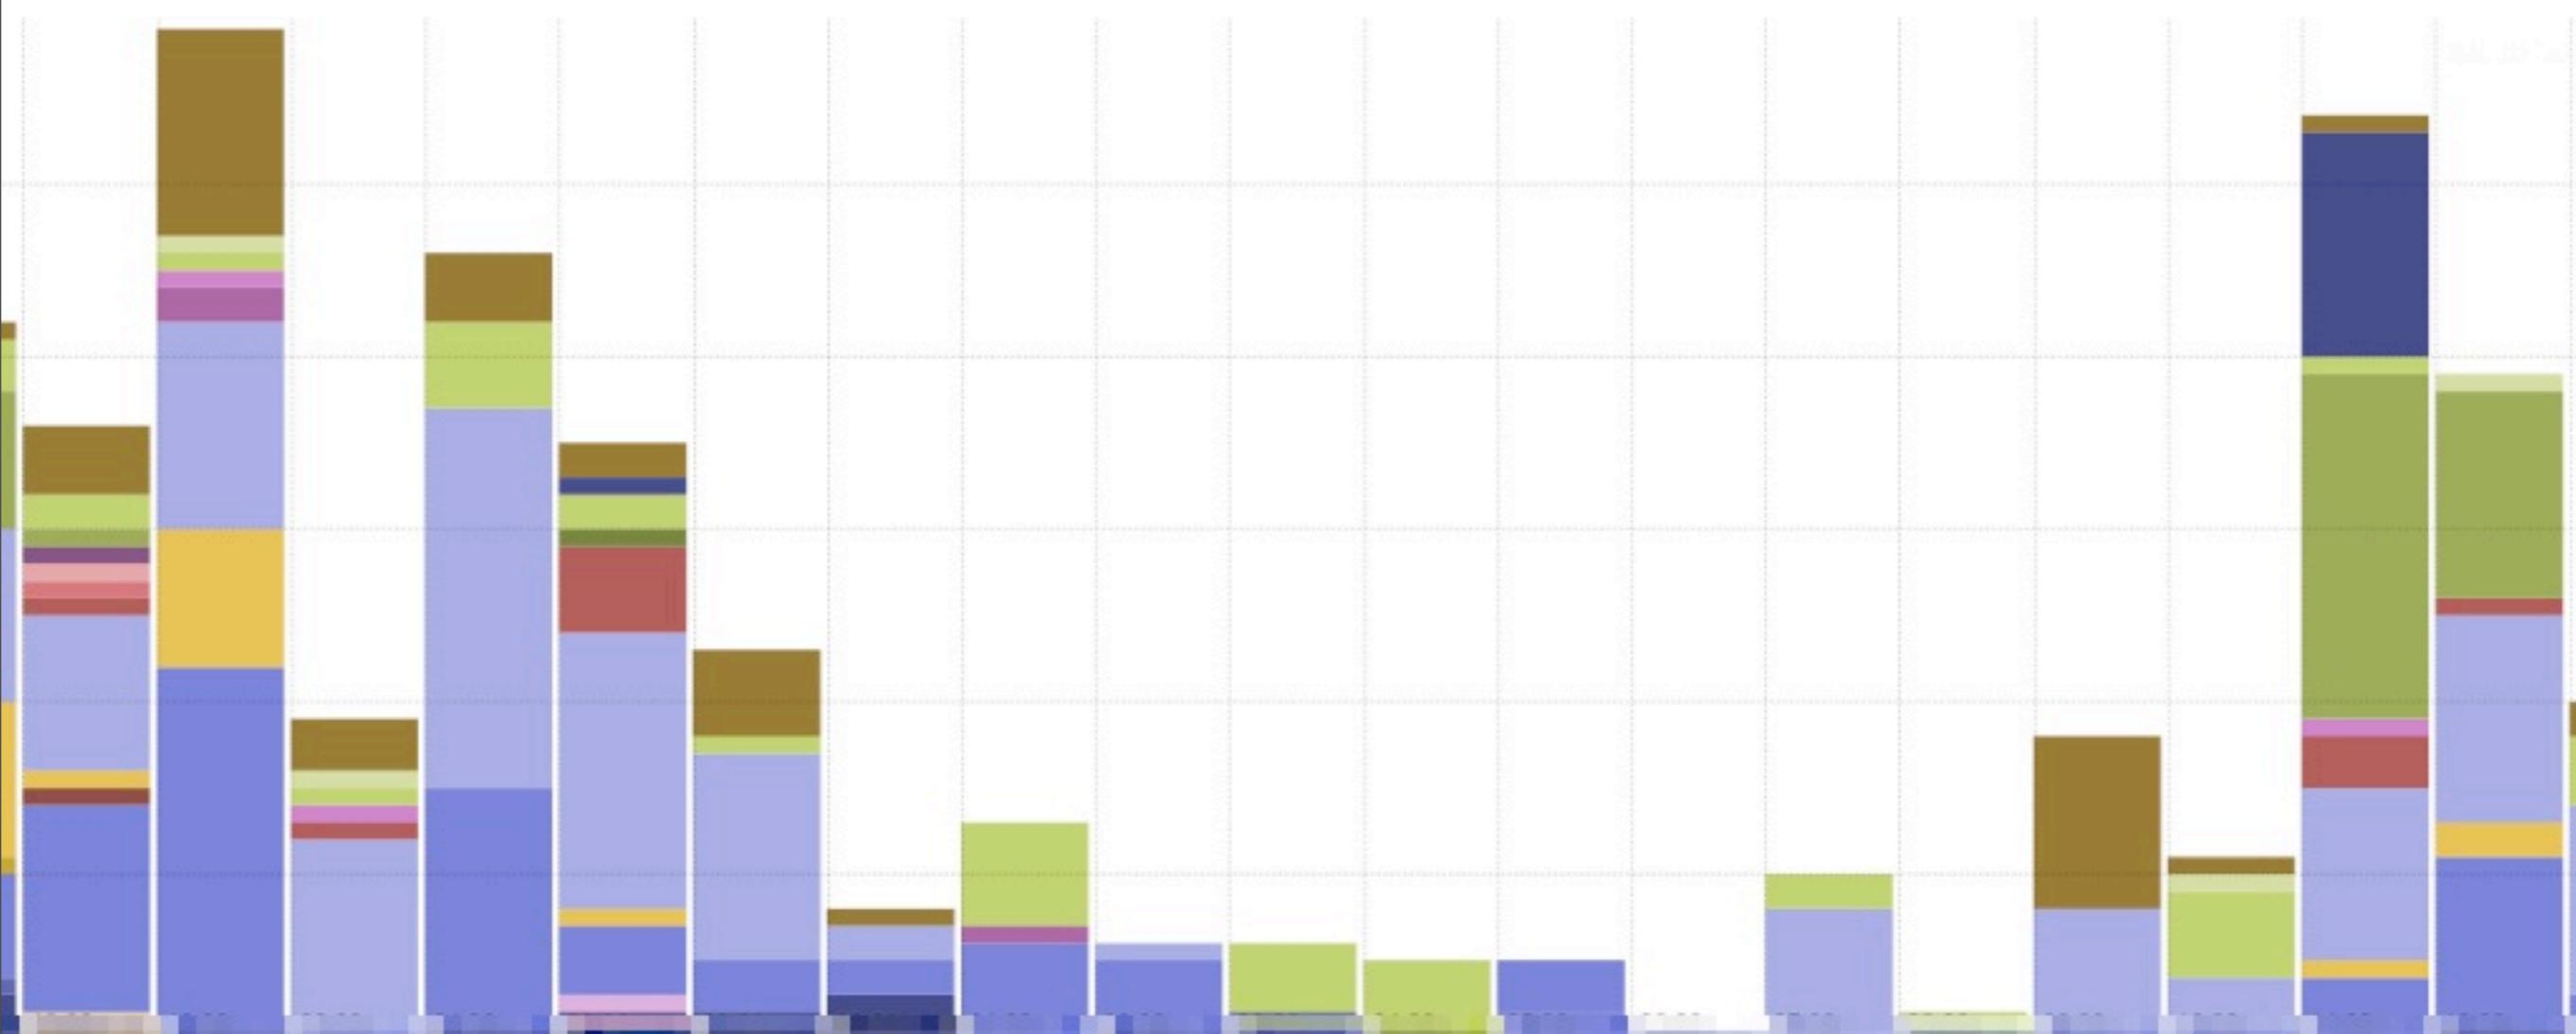
PROFIT!

Also, we know that we now know
how to identify exceptions that
are 404s in reality

exceptions => 404s

that look like 404s

are actually

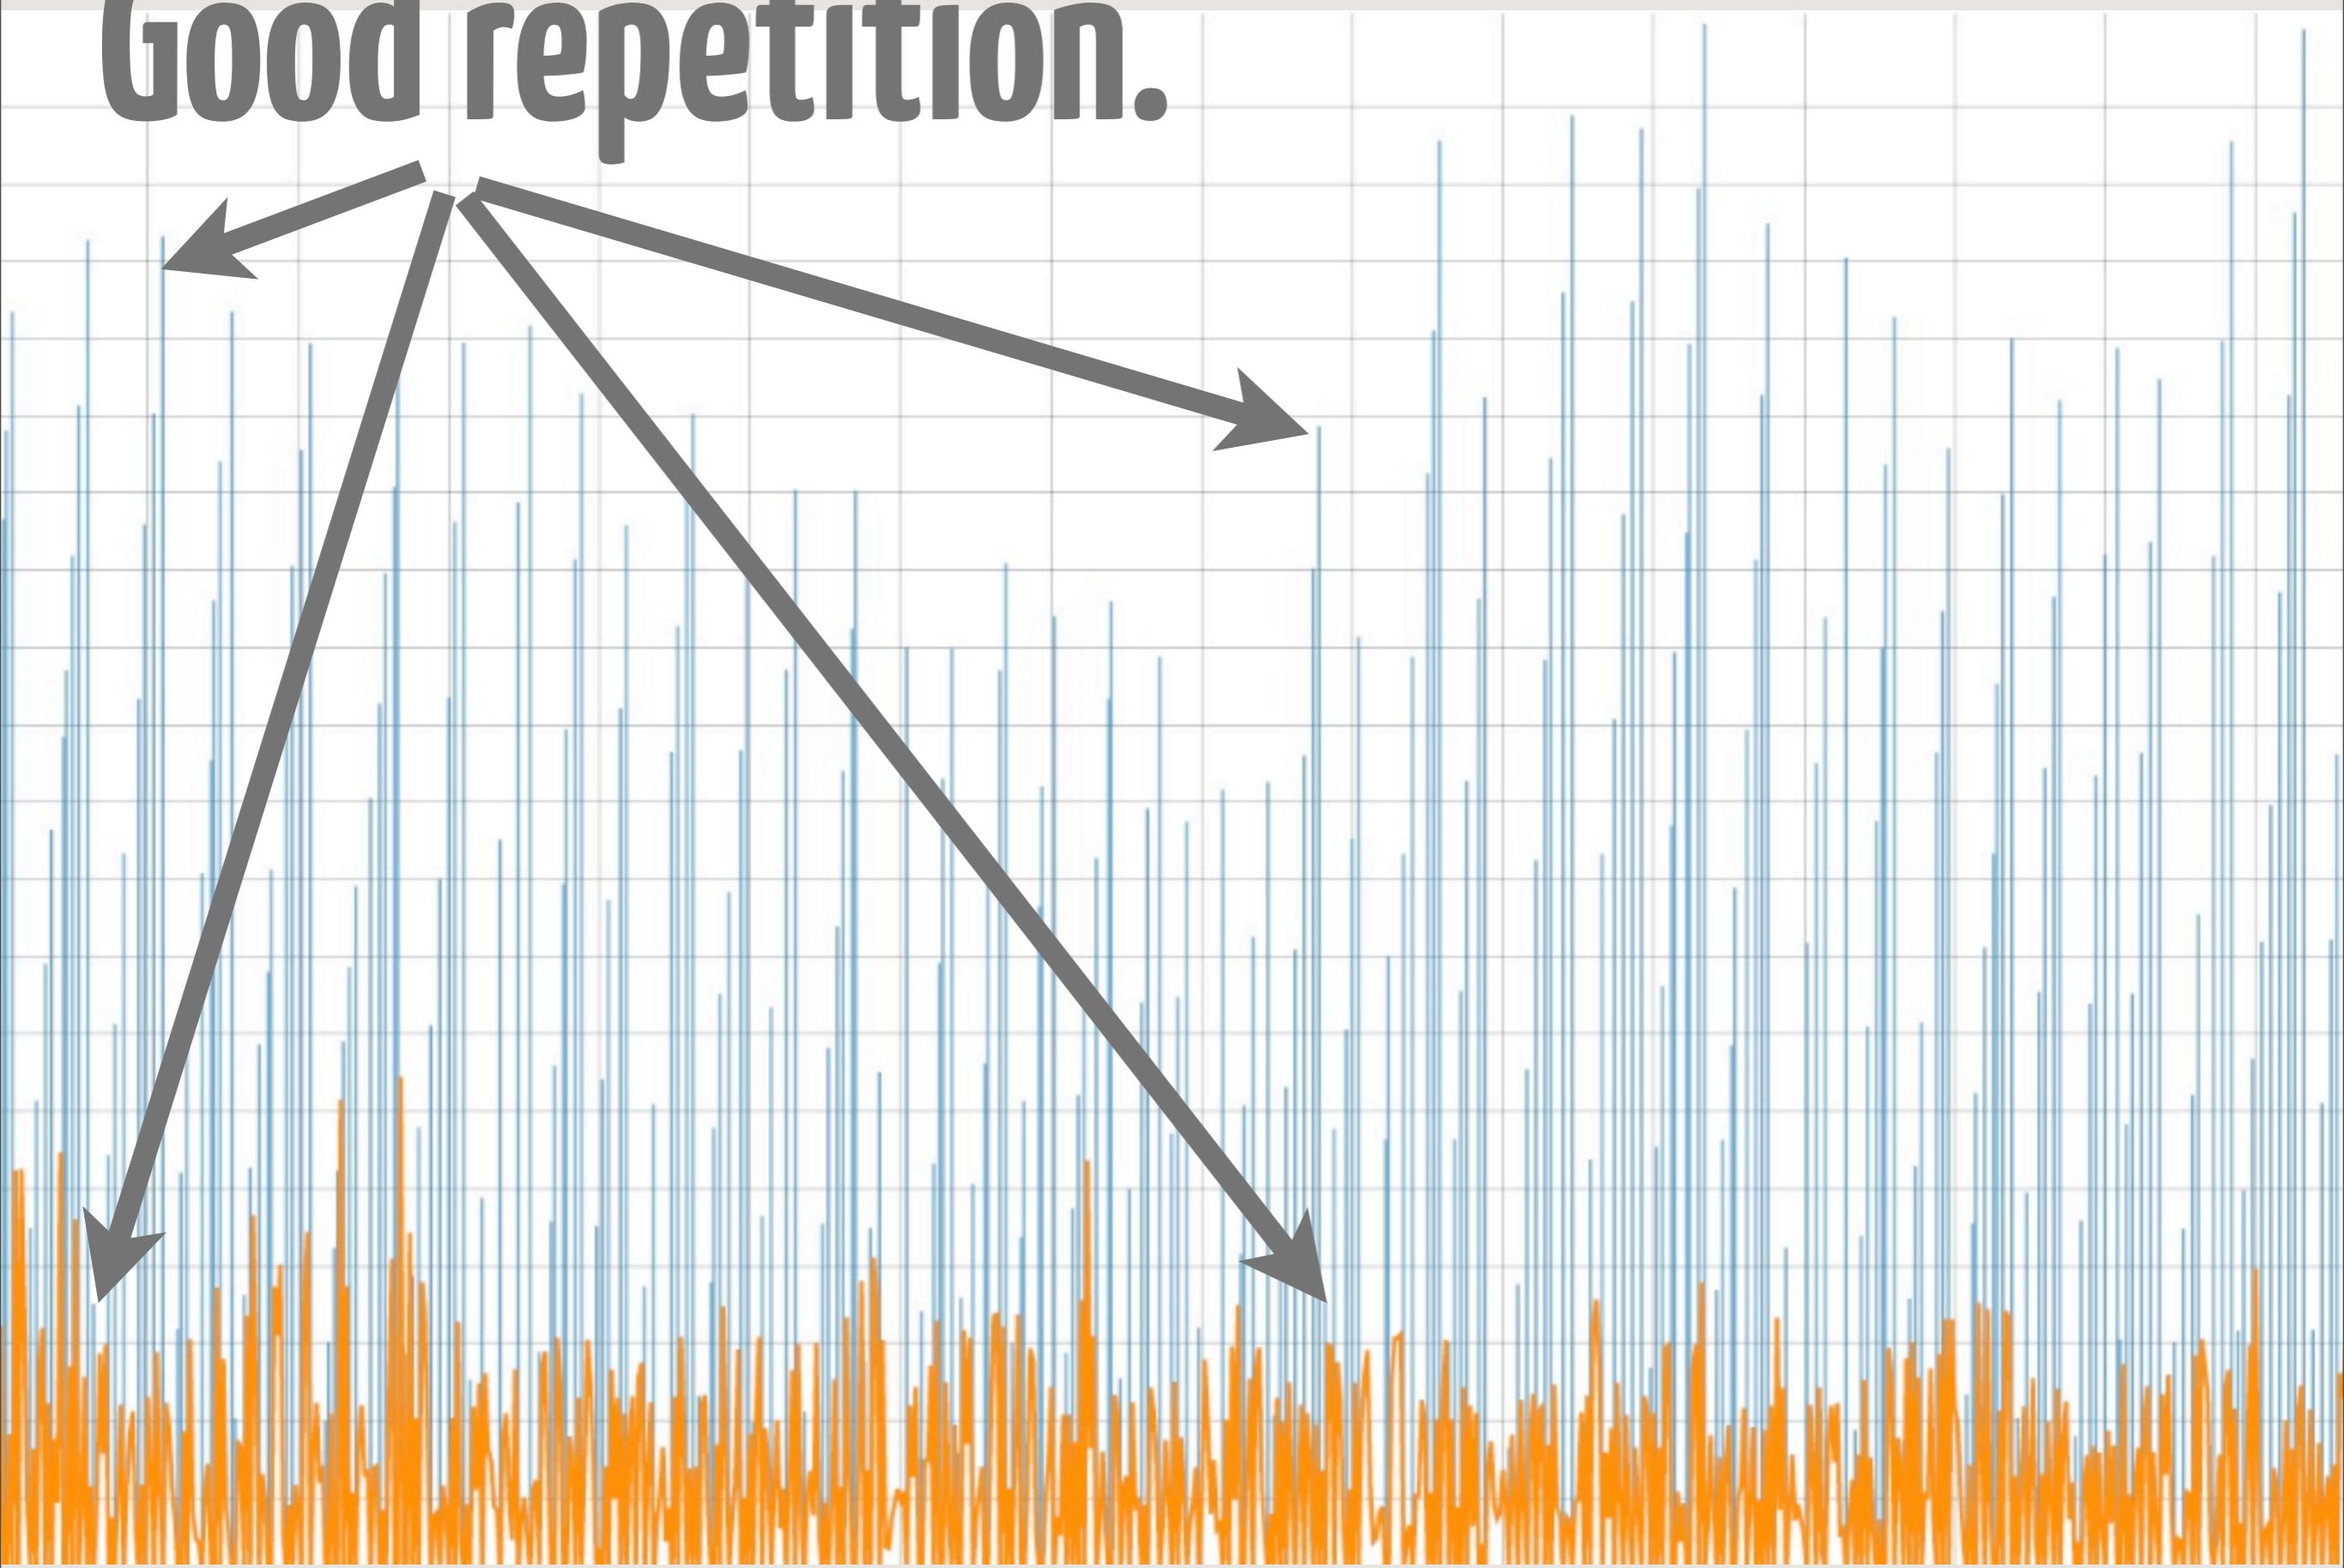


Still some repetition, but much
better

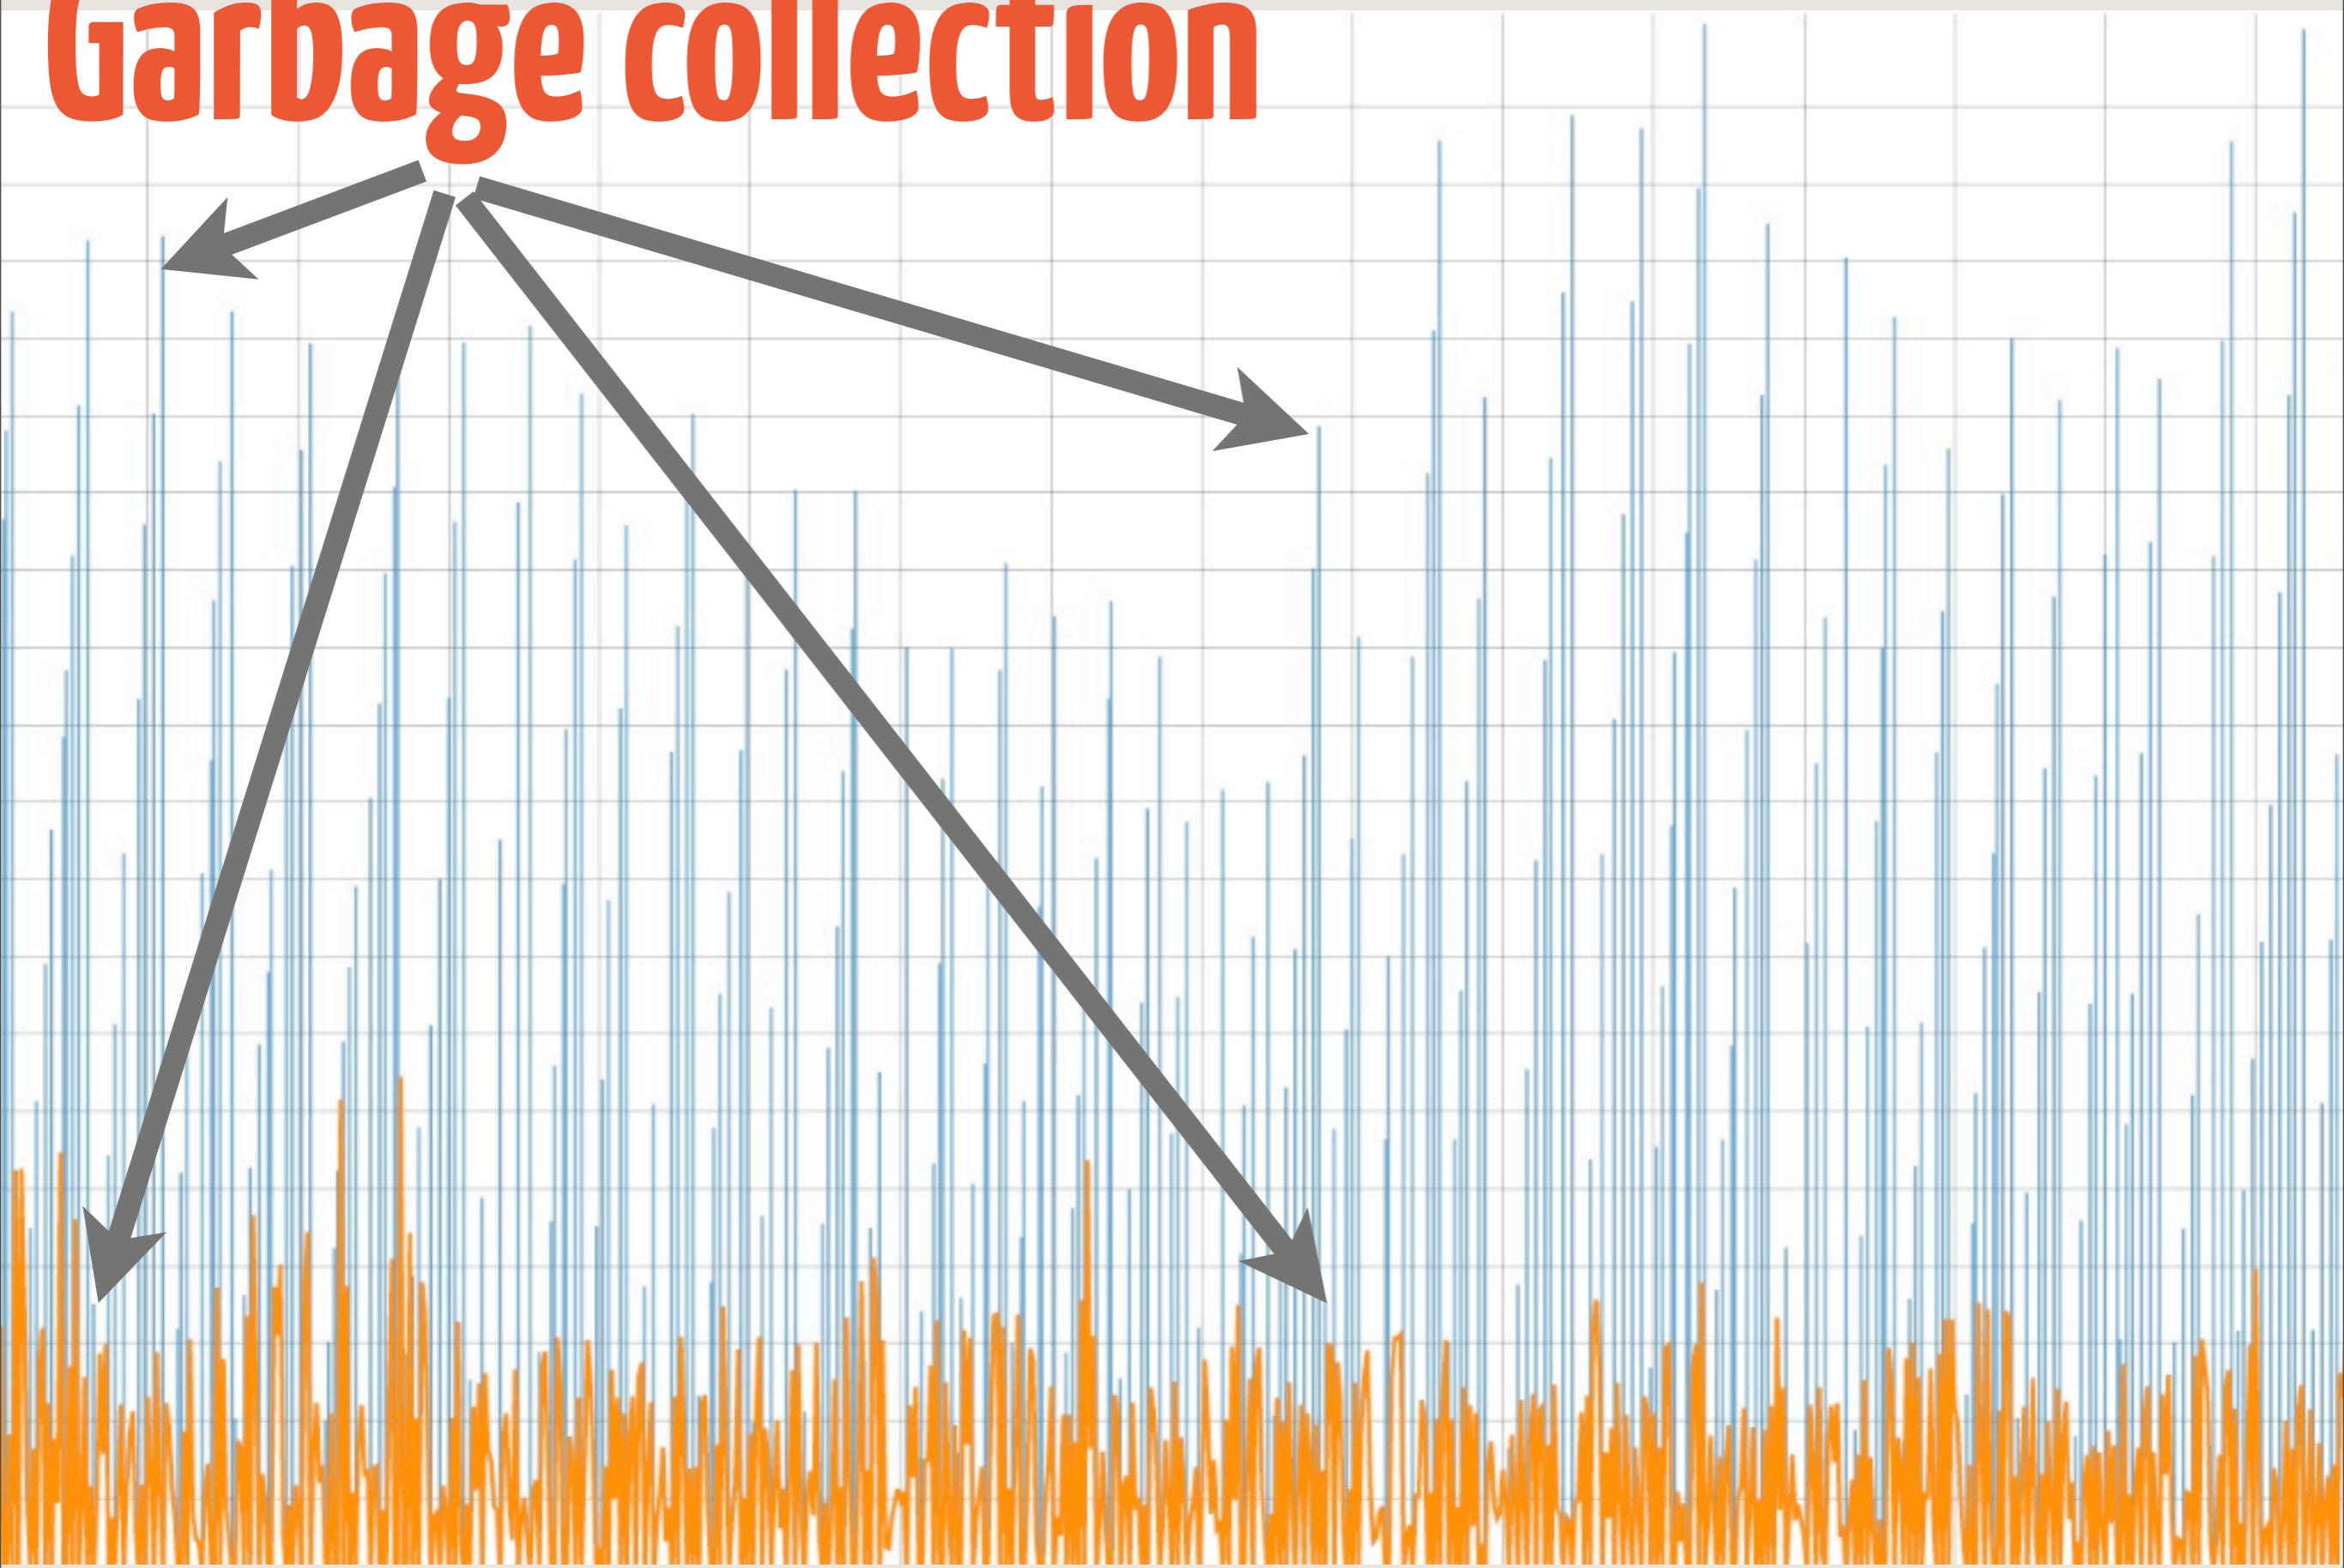
Repetition is *not* necessarily bad

But you must understand **why** it
occurs

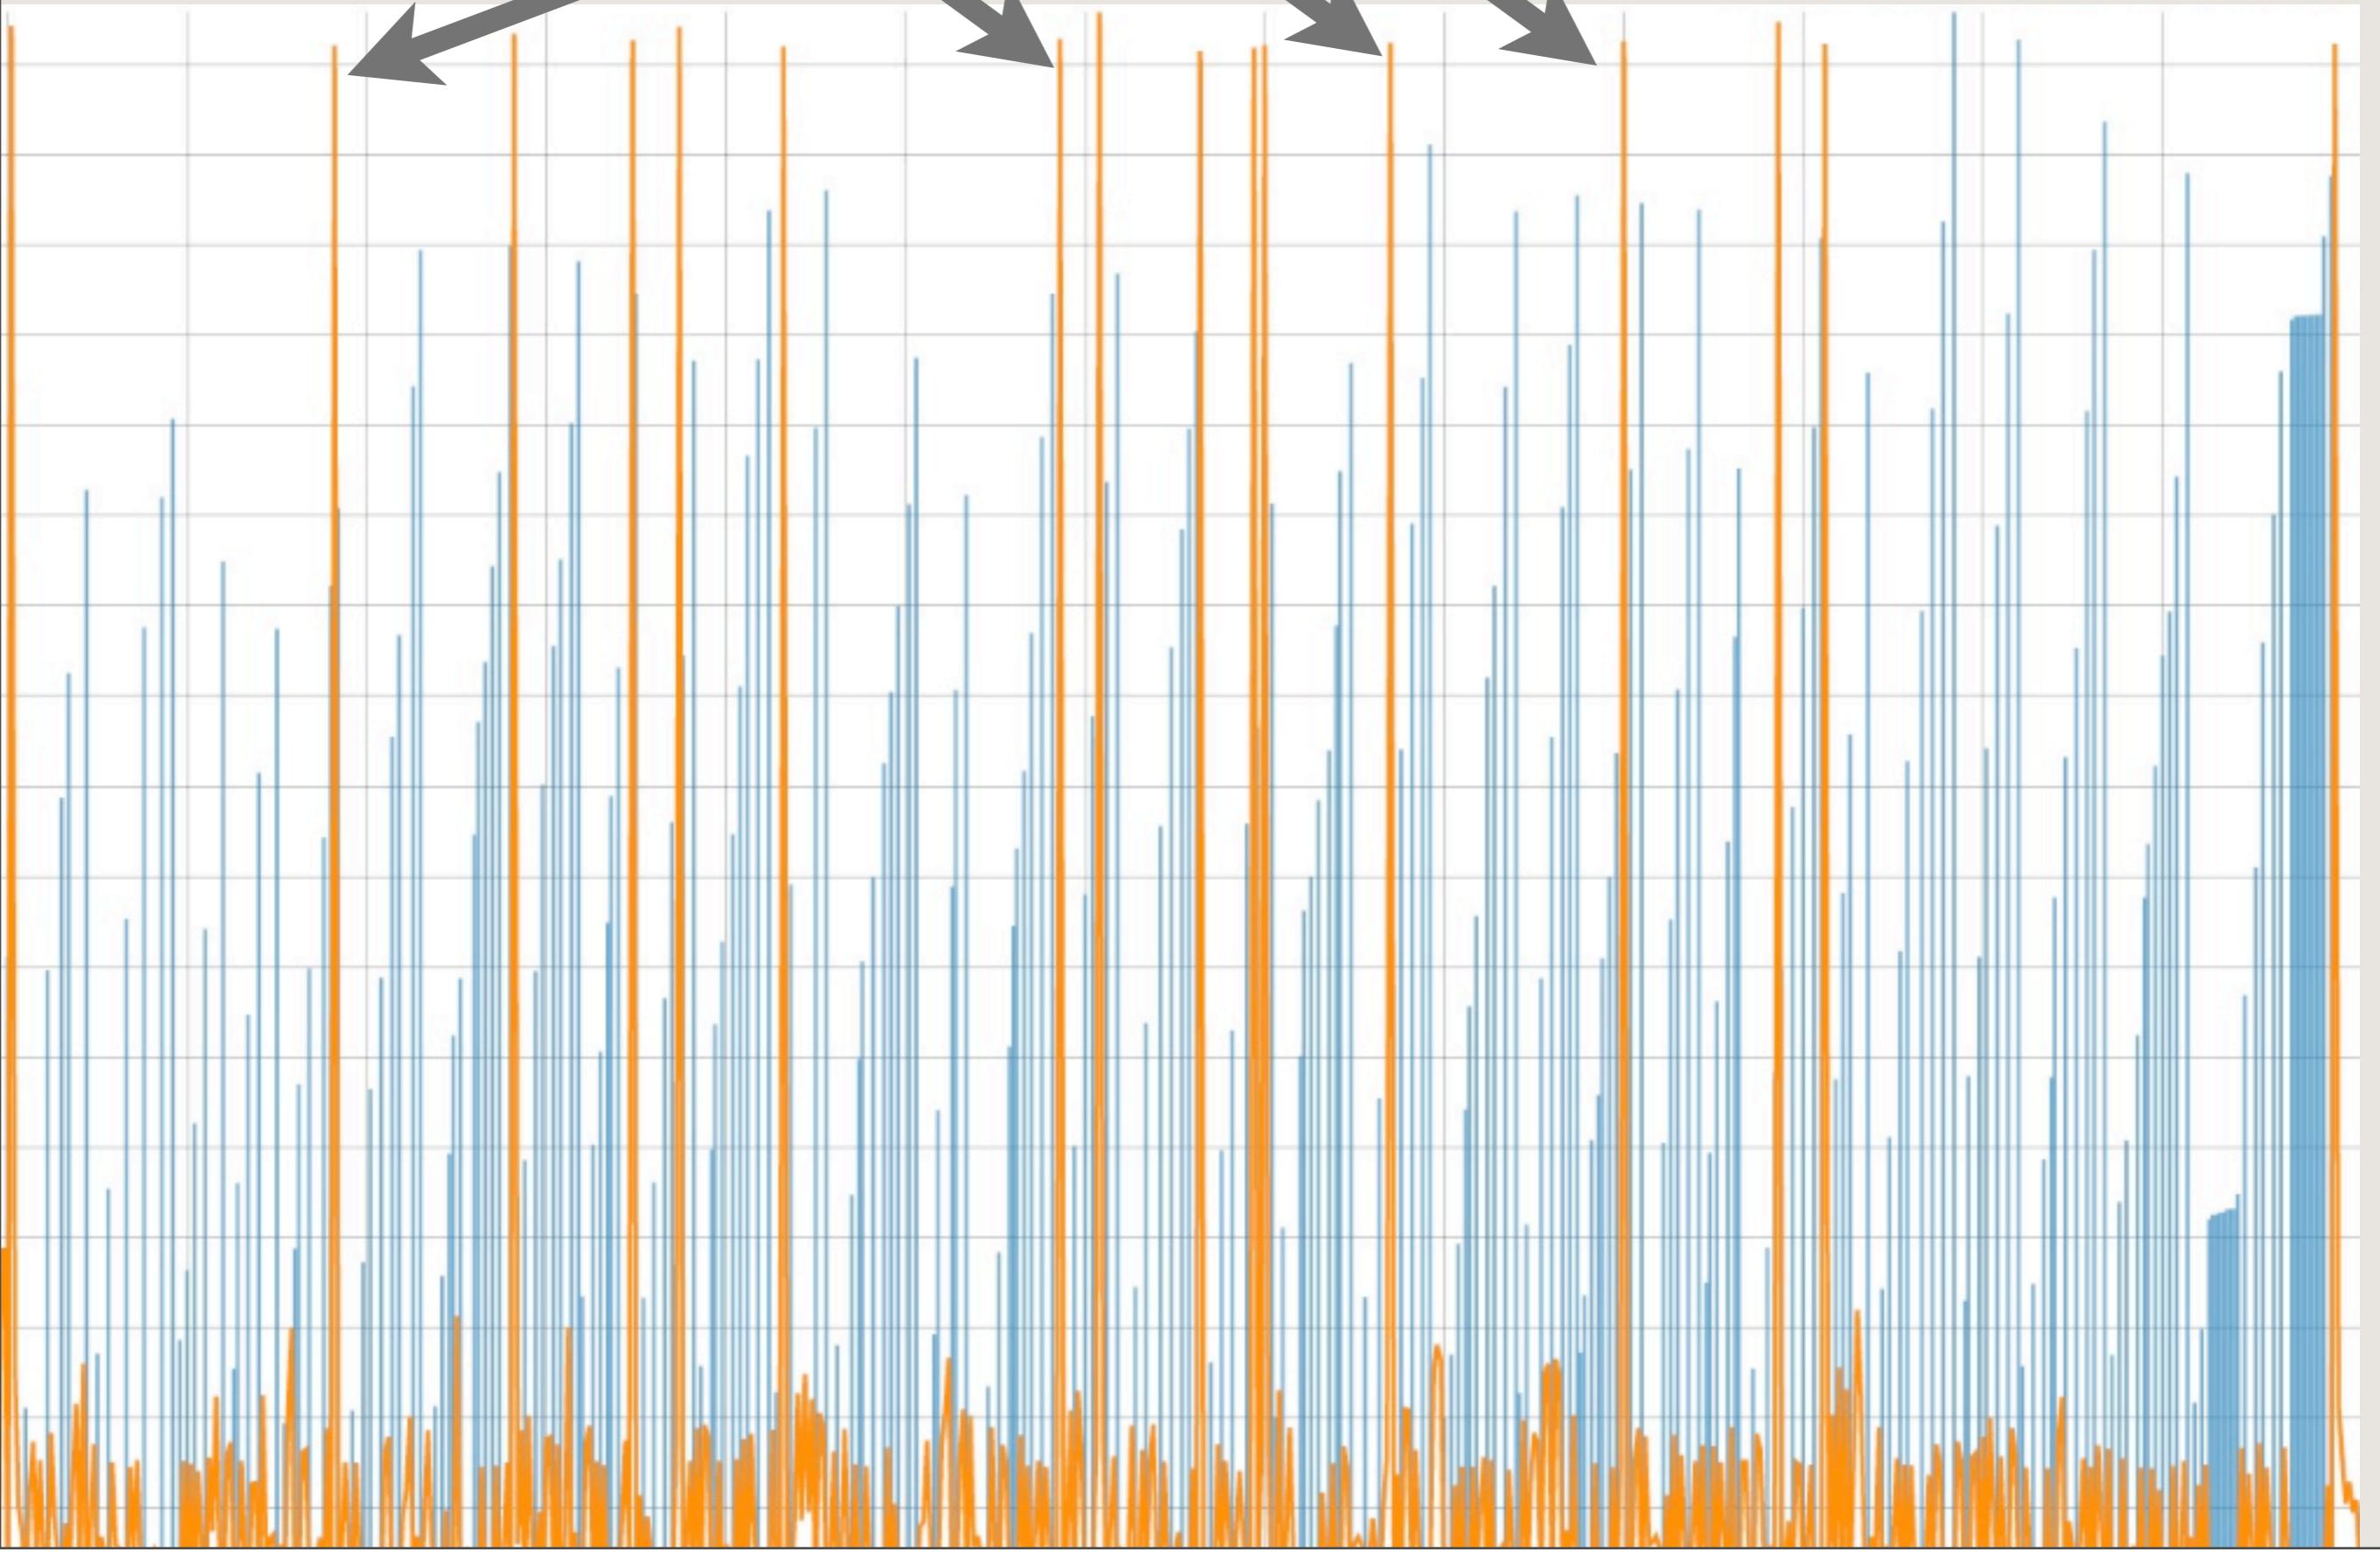
Good repetition.



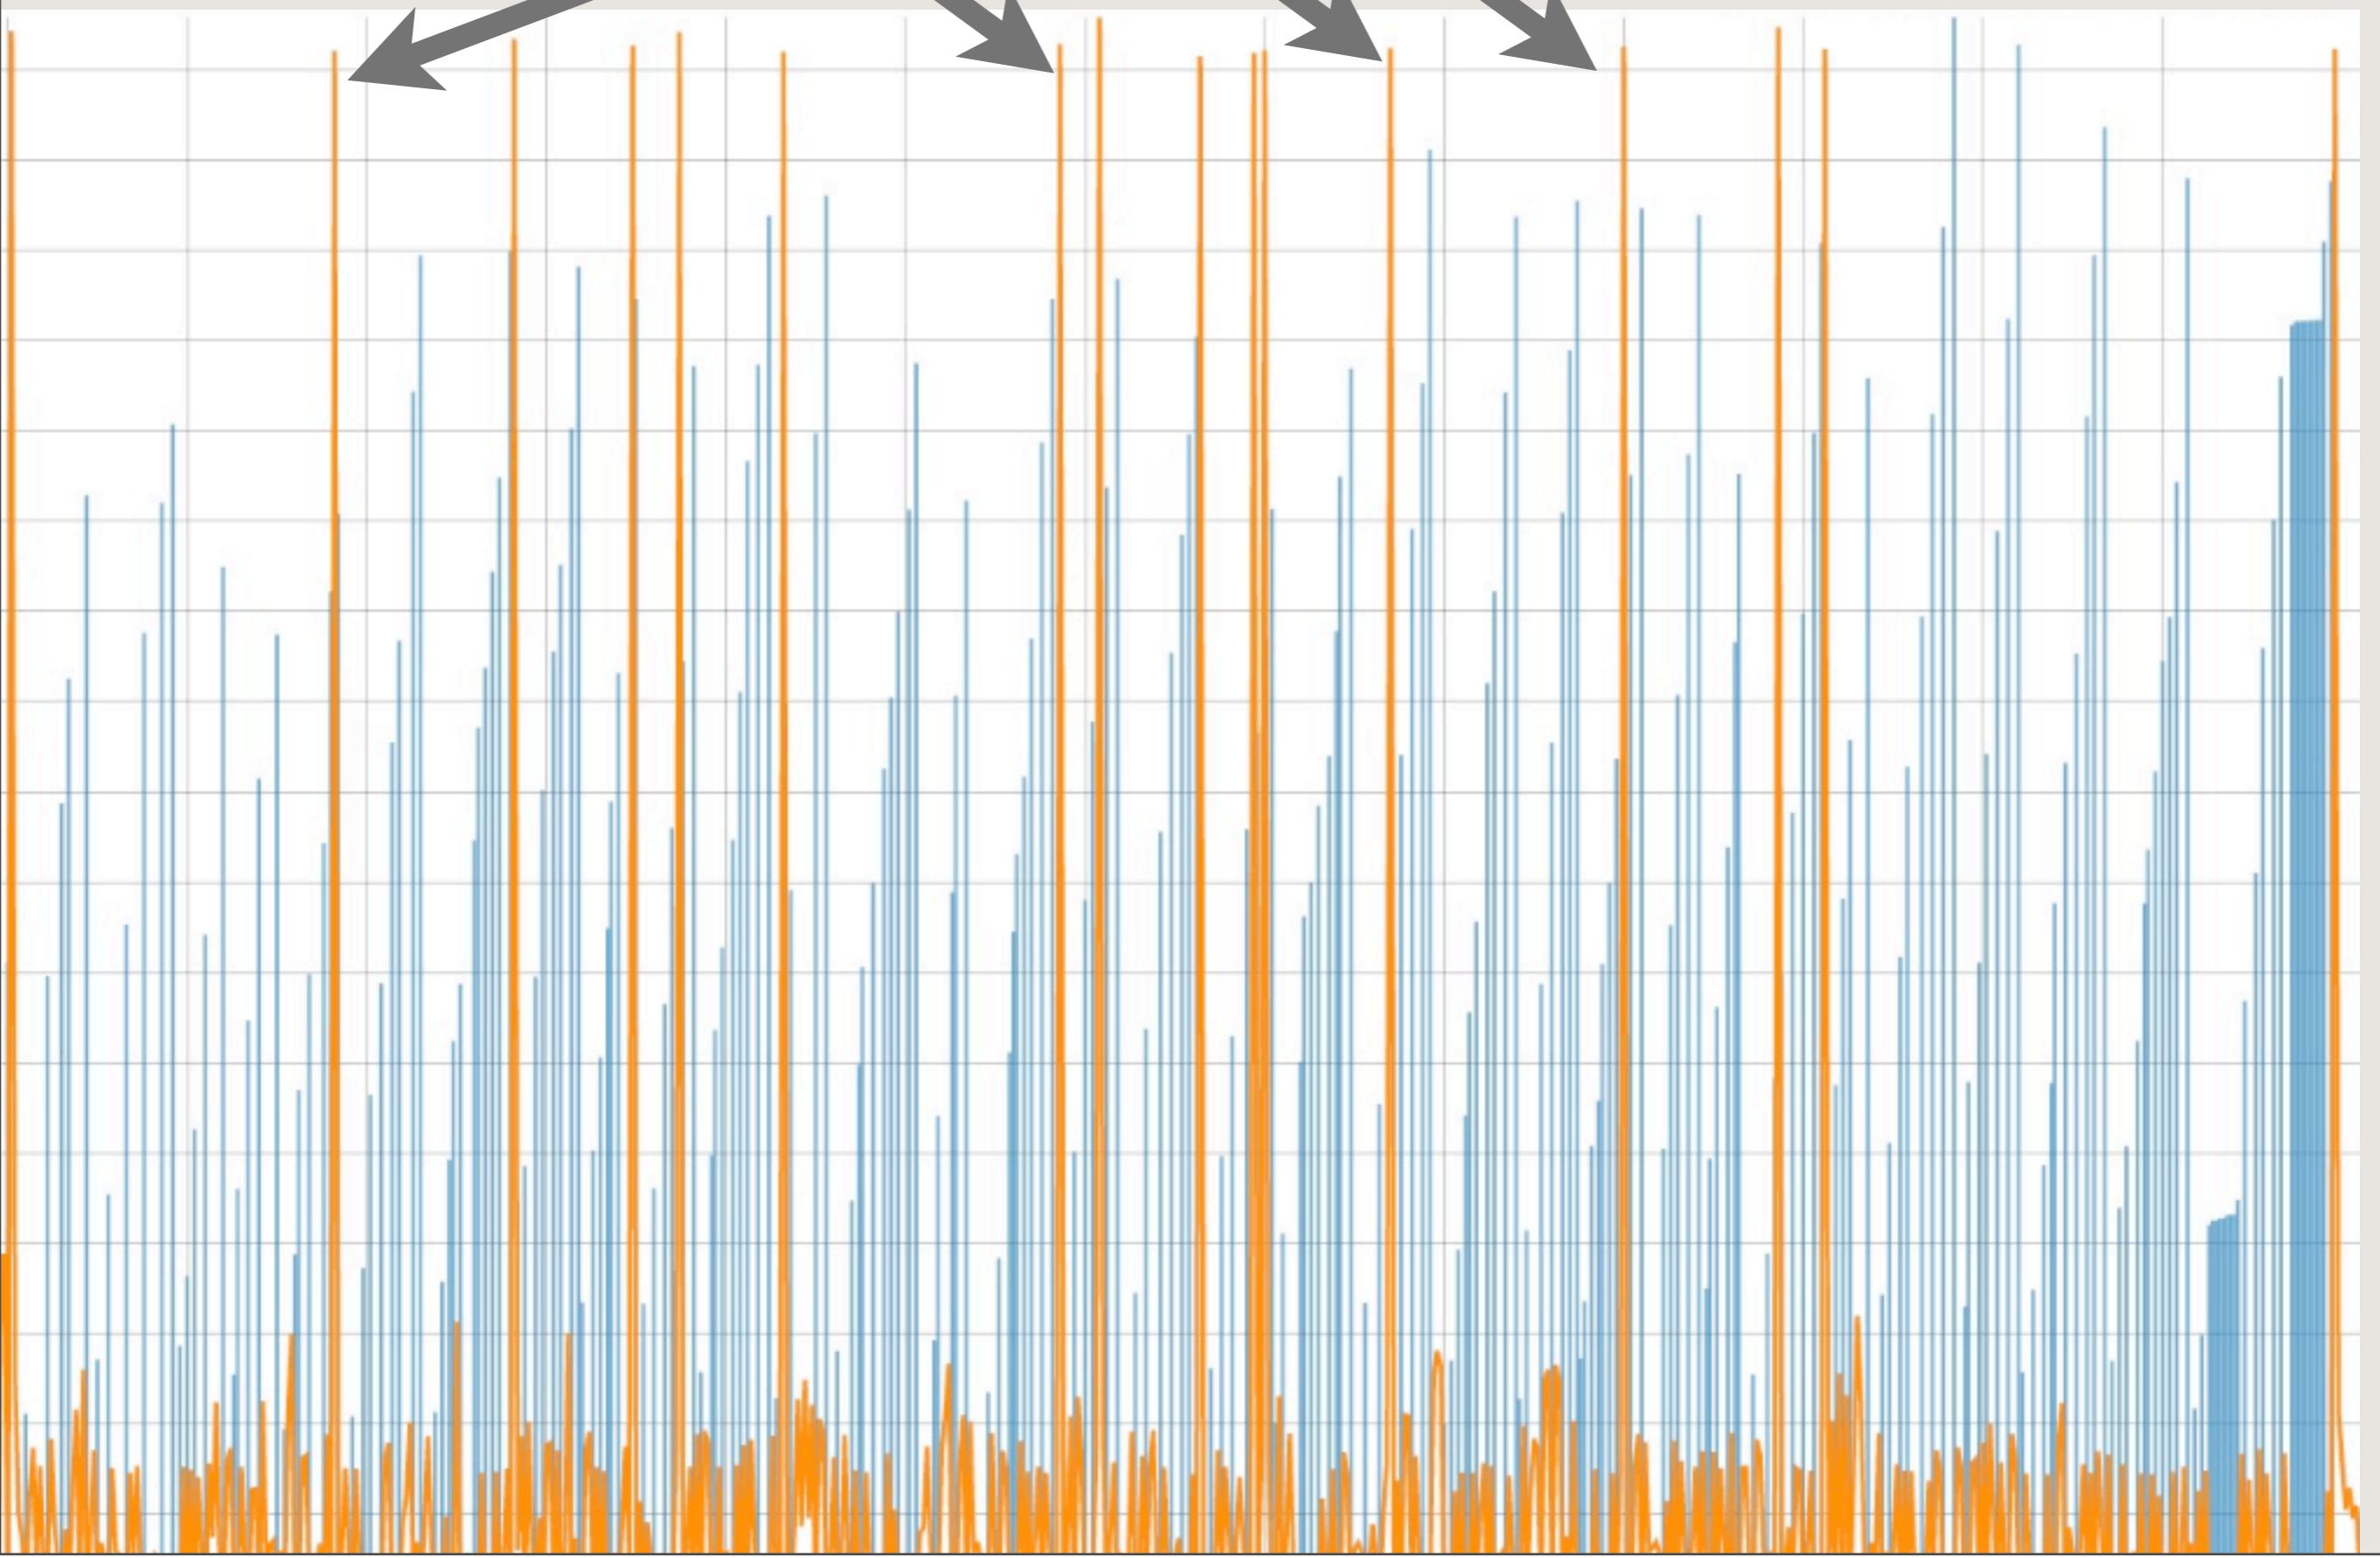
Garbage collection



Bad Repetition

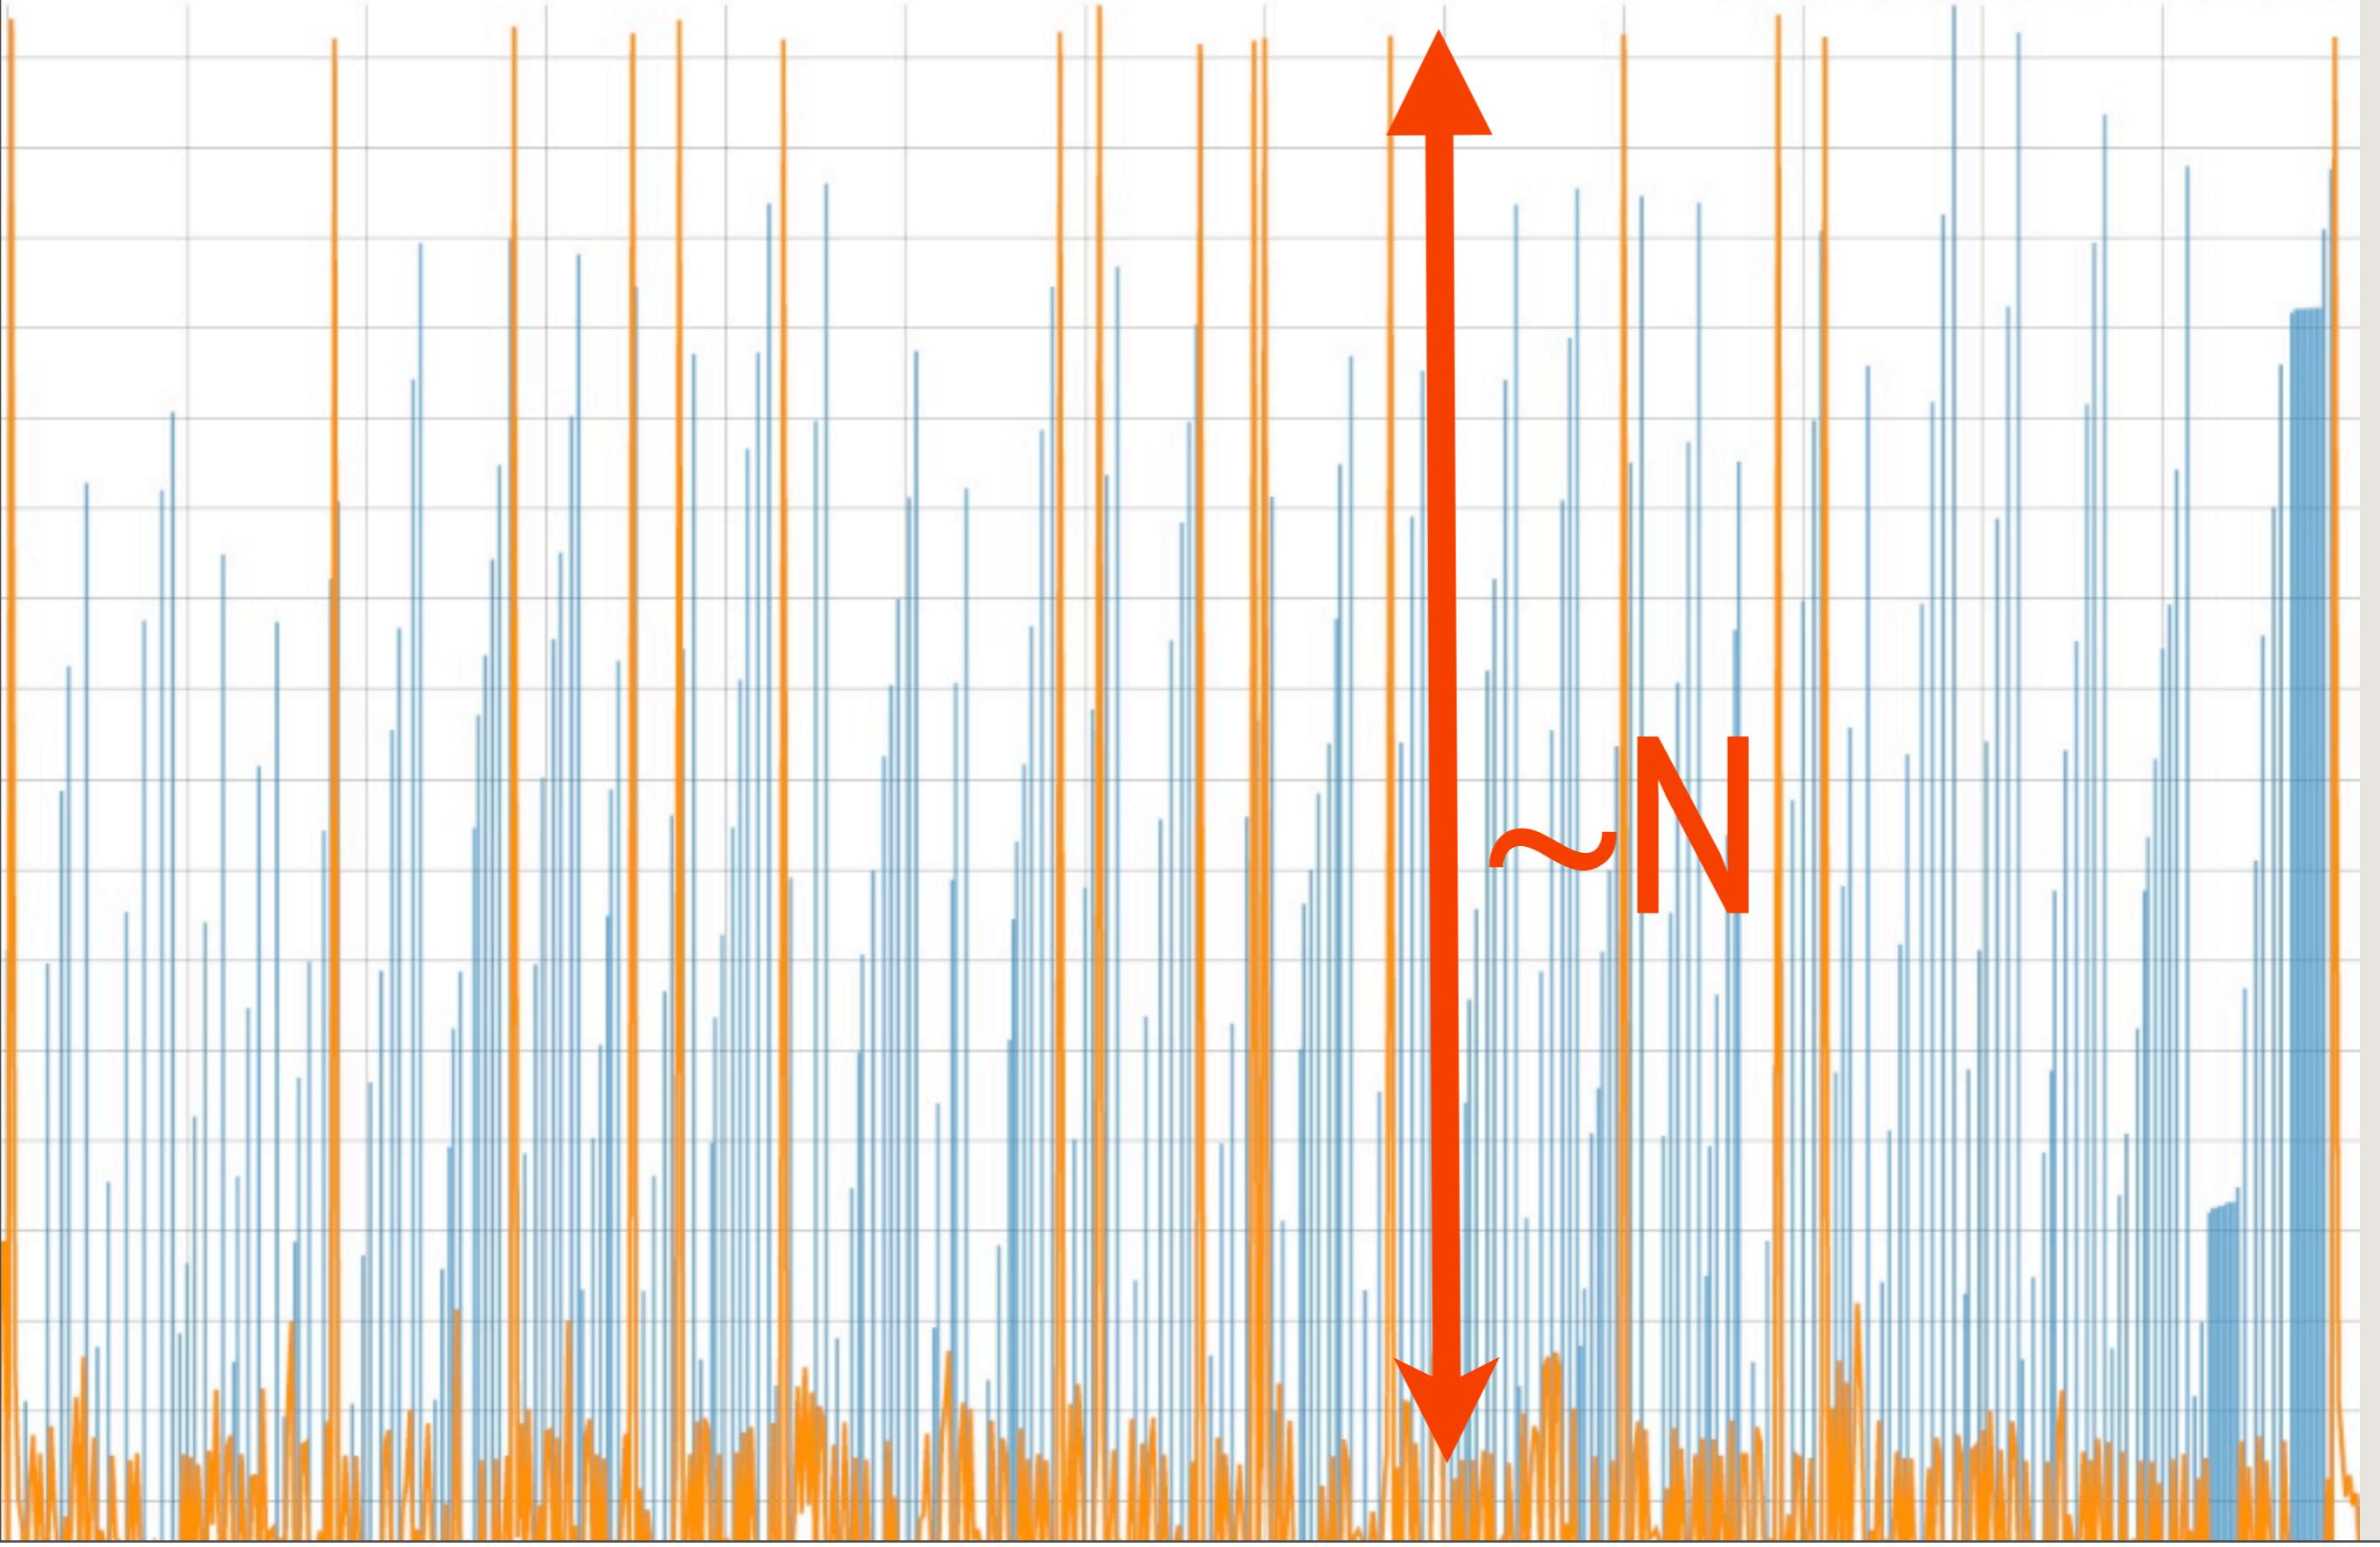


Performance decrease



any other hints?

hmm...
could not obtain database
connection within **N** sec.



tweak parallelism / connection
pool size.

PROFIT!

or...

suddenly, application report **rate**
decreases

server responds **only to every
third/fourth request**

works absolutely fine
after restart

but behavior **recurs** after 5-7
minutes...

somewhere in log files:

Too many open files . . .

and

java.net.UnknownHostException . . .

1.Sof

java	25673	790u	IPv6	53685187	0t0	UDP *:15596
java	25673	791u	IPv6	53687041	0t0	UDP *:42130
java	25673	792u	IPv6	53685196	0t0	UDP *:20910
java	25673	793u	IPv6	53686984	0t0	UDP *:5284
java	25673	794u	IPv6	53686971	0t0	UDP *:4324
java	25673	795u	IPv6	53686990	0t0	UDP *:62078
java	25673	796u	IPv6	53687213	0t0	UDP *:39681
java	25673	797u	IPv6	53687022	0t0	UDP *:38956
java	25673	798u	IPv6	53687057	0t0	UDP *:56857
java	25673	799u	IPv6	53687051	0t0	UDP *:40148
java	25673	800u	IPv6	53687224	0t0	UDP *:10271
java	25673	801u	IPv6	53687163	0t0	UDP *:46985
java	25673	802u	IPv6	53687197	0t0	UDP *:41271
java	25673	803u	IPv6	53687094	0t0	UDP *:26566
java	25673	804u	IPv6	53687183	0t0	UDP *:45242
java	25673	805u	IPv6	53687170	0t0	UDP *:8359
java	25673	806u	IPv6	53687206	0t0	UDP *:41881
java	25673	807u	IPv6	53687234	0t0	UDP *:20064
java	25673	808u	IPv6	53687232	0t0	UDP *:27985
java	25673	809u	IPv6	53687222	0t0	UDP *:23808
java	25673	810u	IPv6	53687241	0t0	UDP *:1723
java	25673	811u	IPv6	53688629	0t0	UDP *:21399
java	25673	812u	IPv6	53687243	0t0	UDP *:12013
java	25673	813u	IPv6	53687245	0t0	UDP *:50506
java	25673	814u	IPv6	53687247	0t0	UDP *:38468
java	25673	815u	IPv6	53687249	0t0	UDP *:46785
java	25673	816u	IPv6	53687251	0t0	UDP *:11629
java	25673	817u	IPv6	53687254	0t0	UDP *:4663
java	25673	818u	IPv6	53687256	0t0	UDP *:5736
java	25673	819u	IPv6	53687258	0t0	UDP *:60244
java	25673	820u	IPv6	53687260	0t0	UDP *:15657
java	25673	821u	IPv6	53687262	0t0	UDP *:34269
java	25673	822u	IPv6	53687264	0t0	UDP *:10233
java	25673	823u	IPv6	53687266	0t0	UDP *:12452
java	25673	824u	IPv6	53687268	0t0	UDP *:18788
java	25673	825u	IPv6	53687286	0t0	UDP *:7832
java	25673	826u	IPv6	53687288	0t0	UDP *:58249
java	25673	827u	IPv6	53687291	0t0	UDP *:39332
java	25673	828u	IPv6	53687293	0t0	UDP *:38187
java	25673	829u	IPv6	53687295	0t0	UDP *:31435
java	25673	830u	IPv6	53687297	0t0	UDP *:34363
java	25673	831u	IPv6	53687299	0t0	UDP *:49518
java	25673	832u	IPv6	53687301	0t0	UDP *:15172
java	25673	833u	IPv6	53687306	0t0	UDP *:17763
java	25673	834u	IPv6	53687308	0t0	UDP *:37129

Occupy UDP or what?

Find connection leak

Fix connection leak

PROFIT!

I would like to

have a production monitoring
system

that will allow me to see these
things early

- very fast and easy access to recent events (recency is arbitrary)
- fast arbitrary queries on recent events
- real-time analytics, report generation
- easy to go back in time
- scalable queries across an entire dataset

- pipelining, ability to inject additional processor
- arbitrary data format
- dumb client
- extensible server (plug-n-play)
- direct access to data, right in my repl

- **easy, extensible HTML / CSS / JS interface**
- **graph generation (real-time, pre-calculated)**
- **websockets for pushing stuff**
- **re-broadcast of incoming events for custom analytics**
- **support for multiple incoming channels**

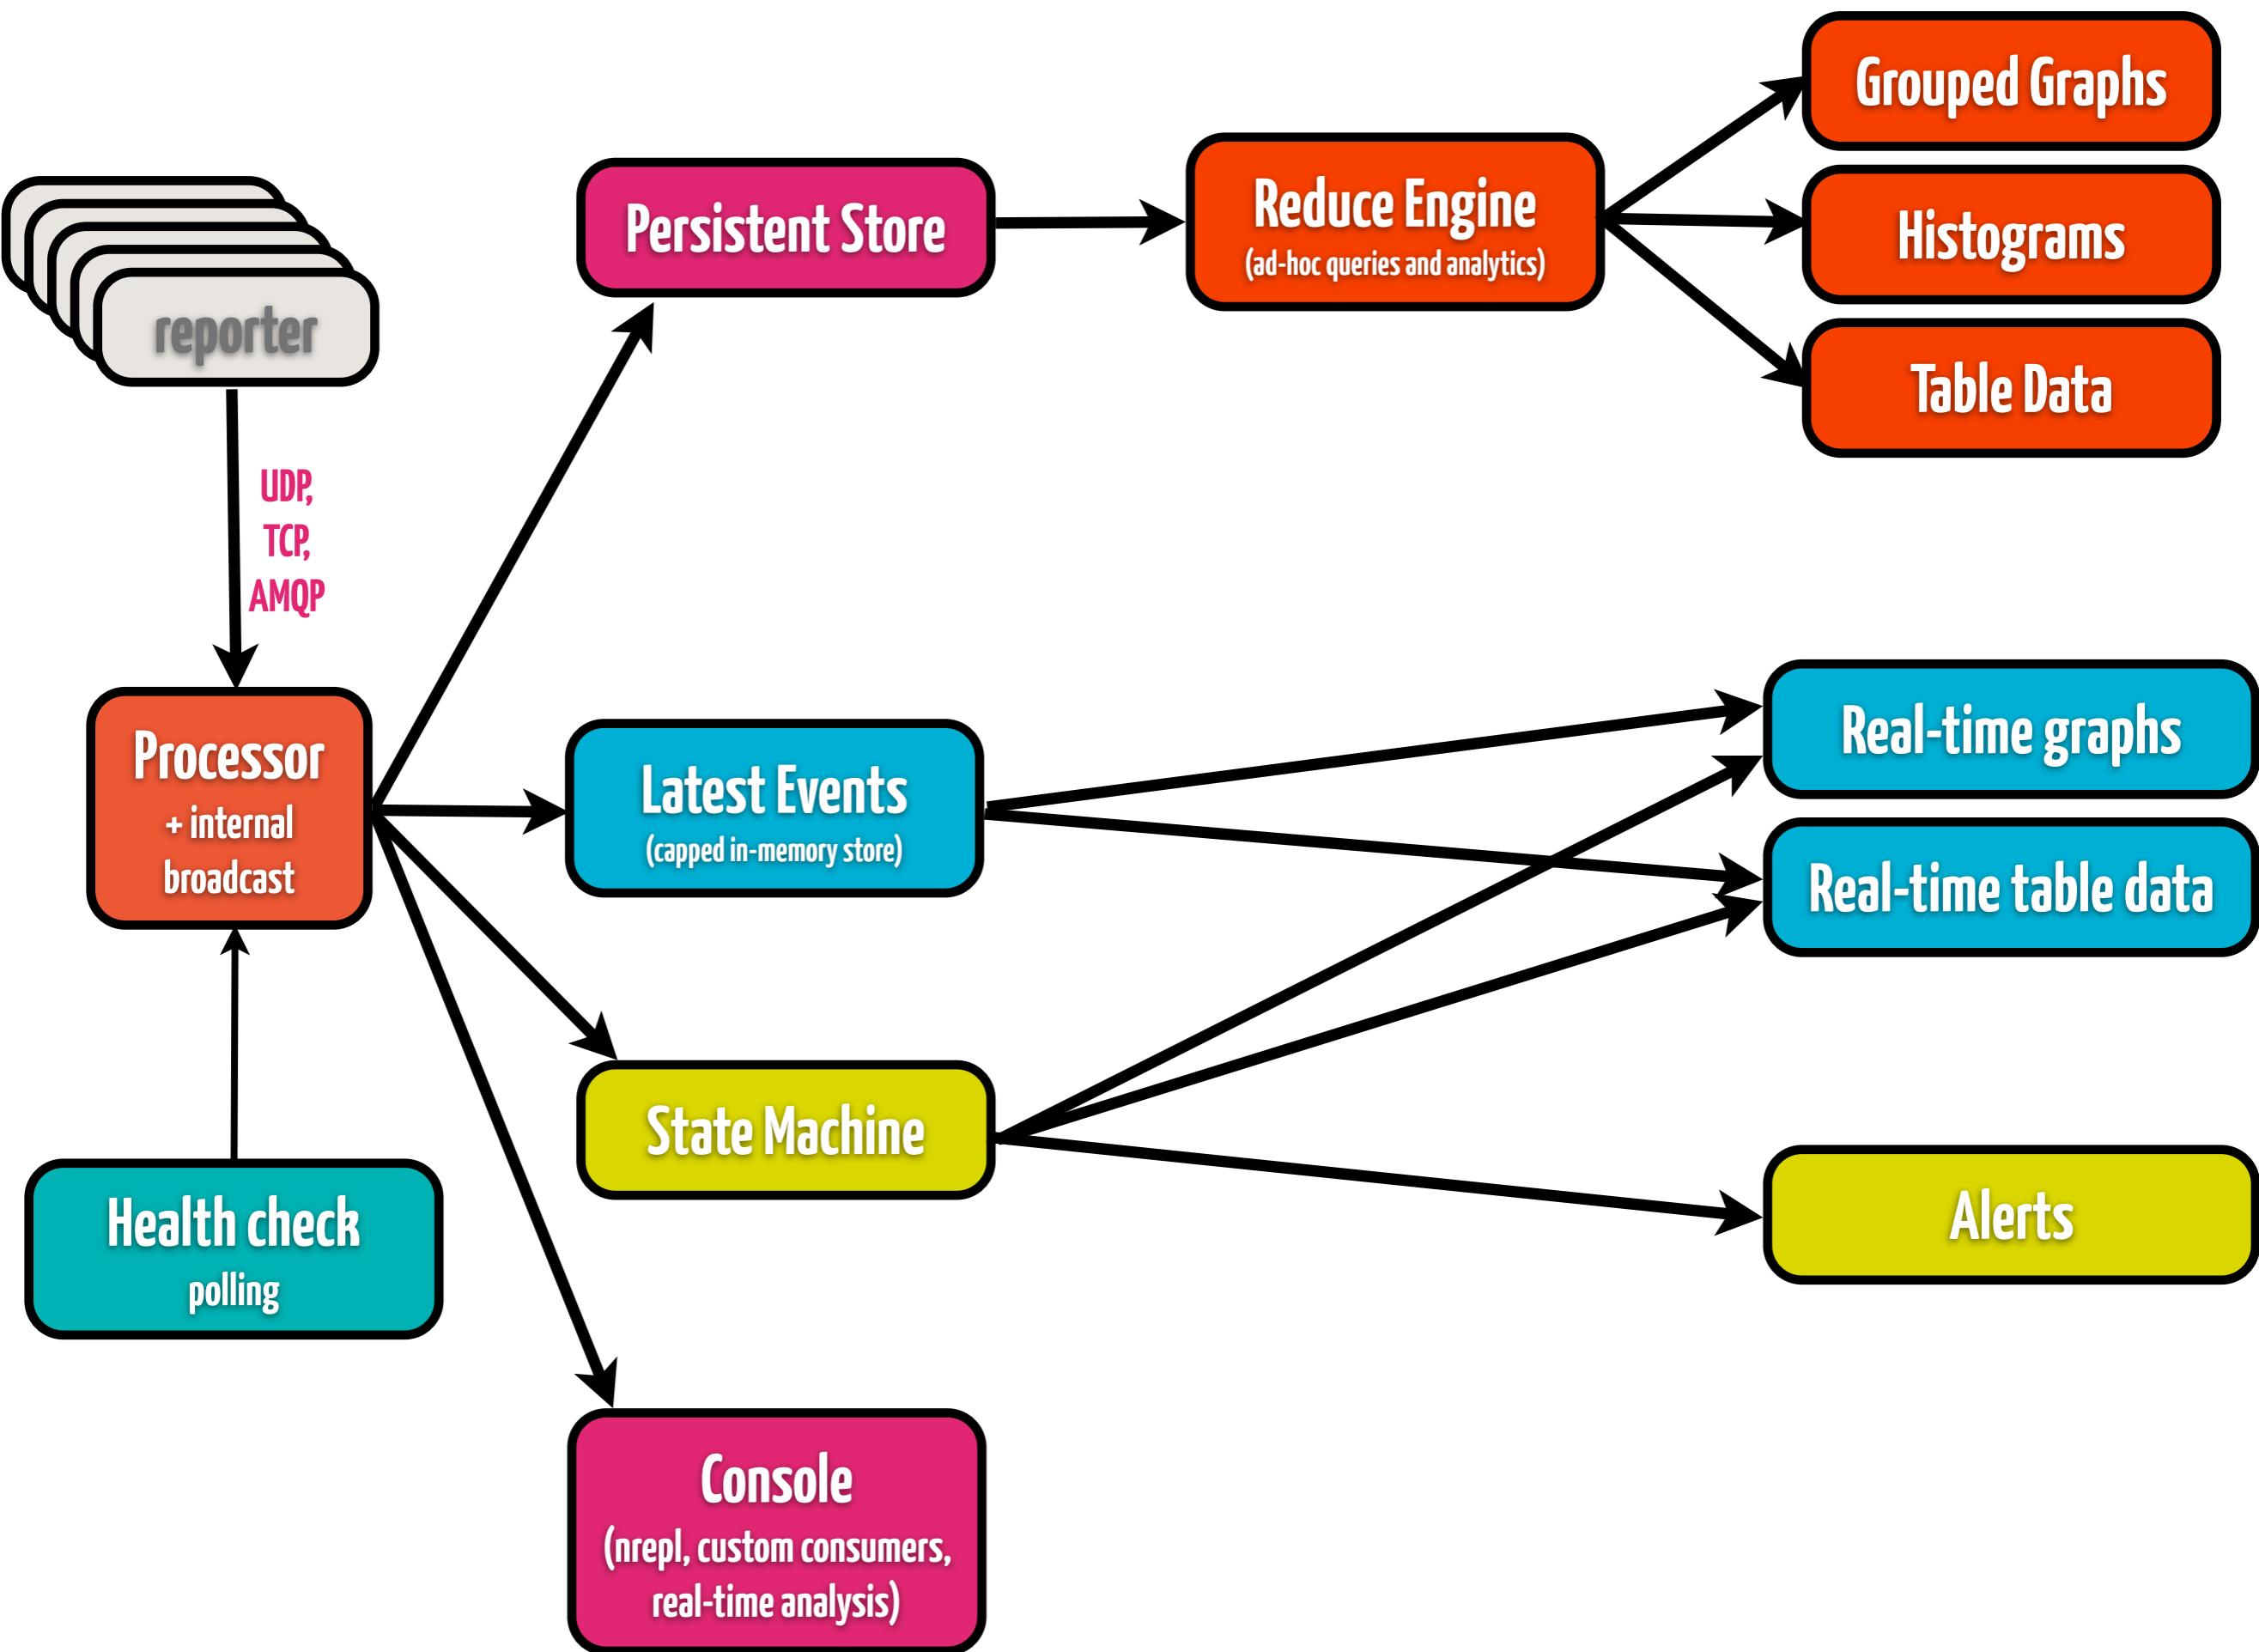
- extensible alerting system (channels)
- extensible alerting system (rules)
- and (probably) many many more, which may be subsets of the mentioned ones in some way

let's draw a simple **schema** that
will represent such a system

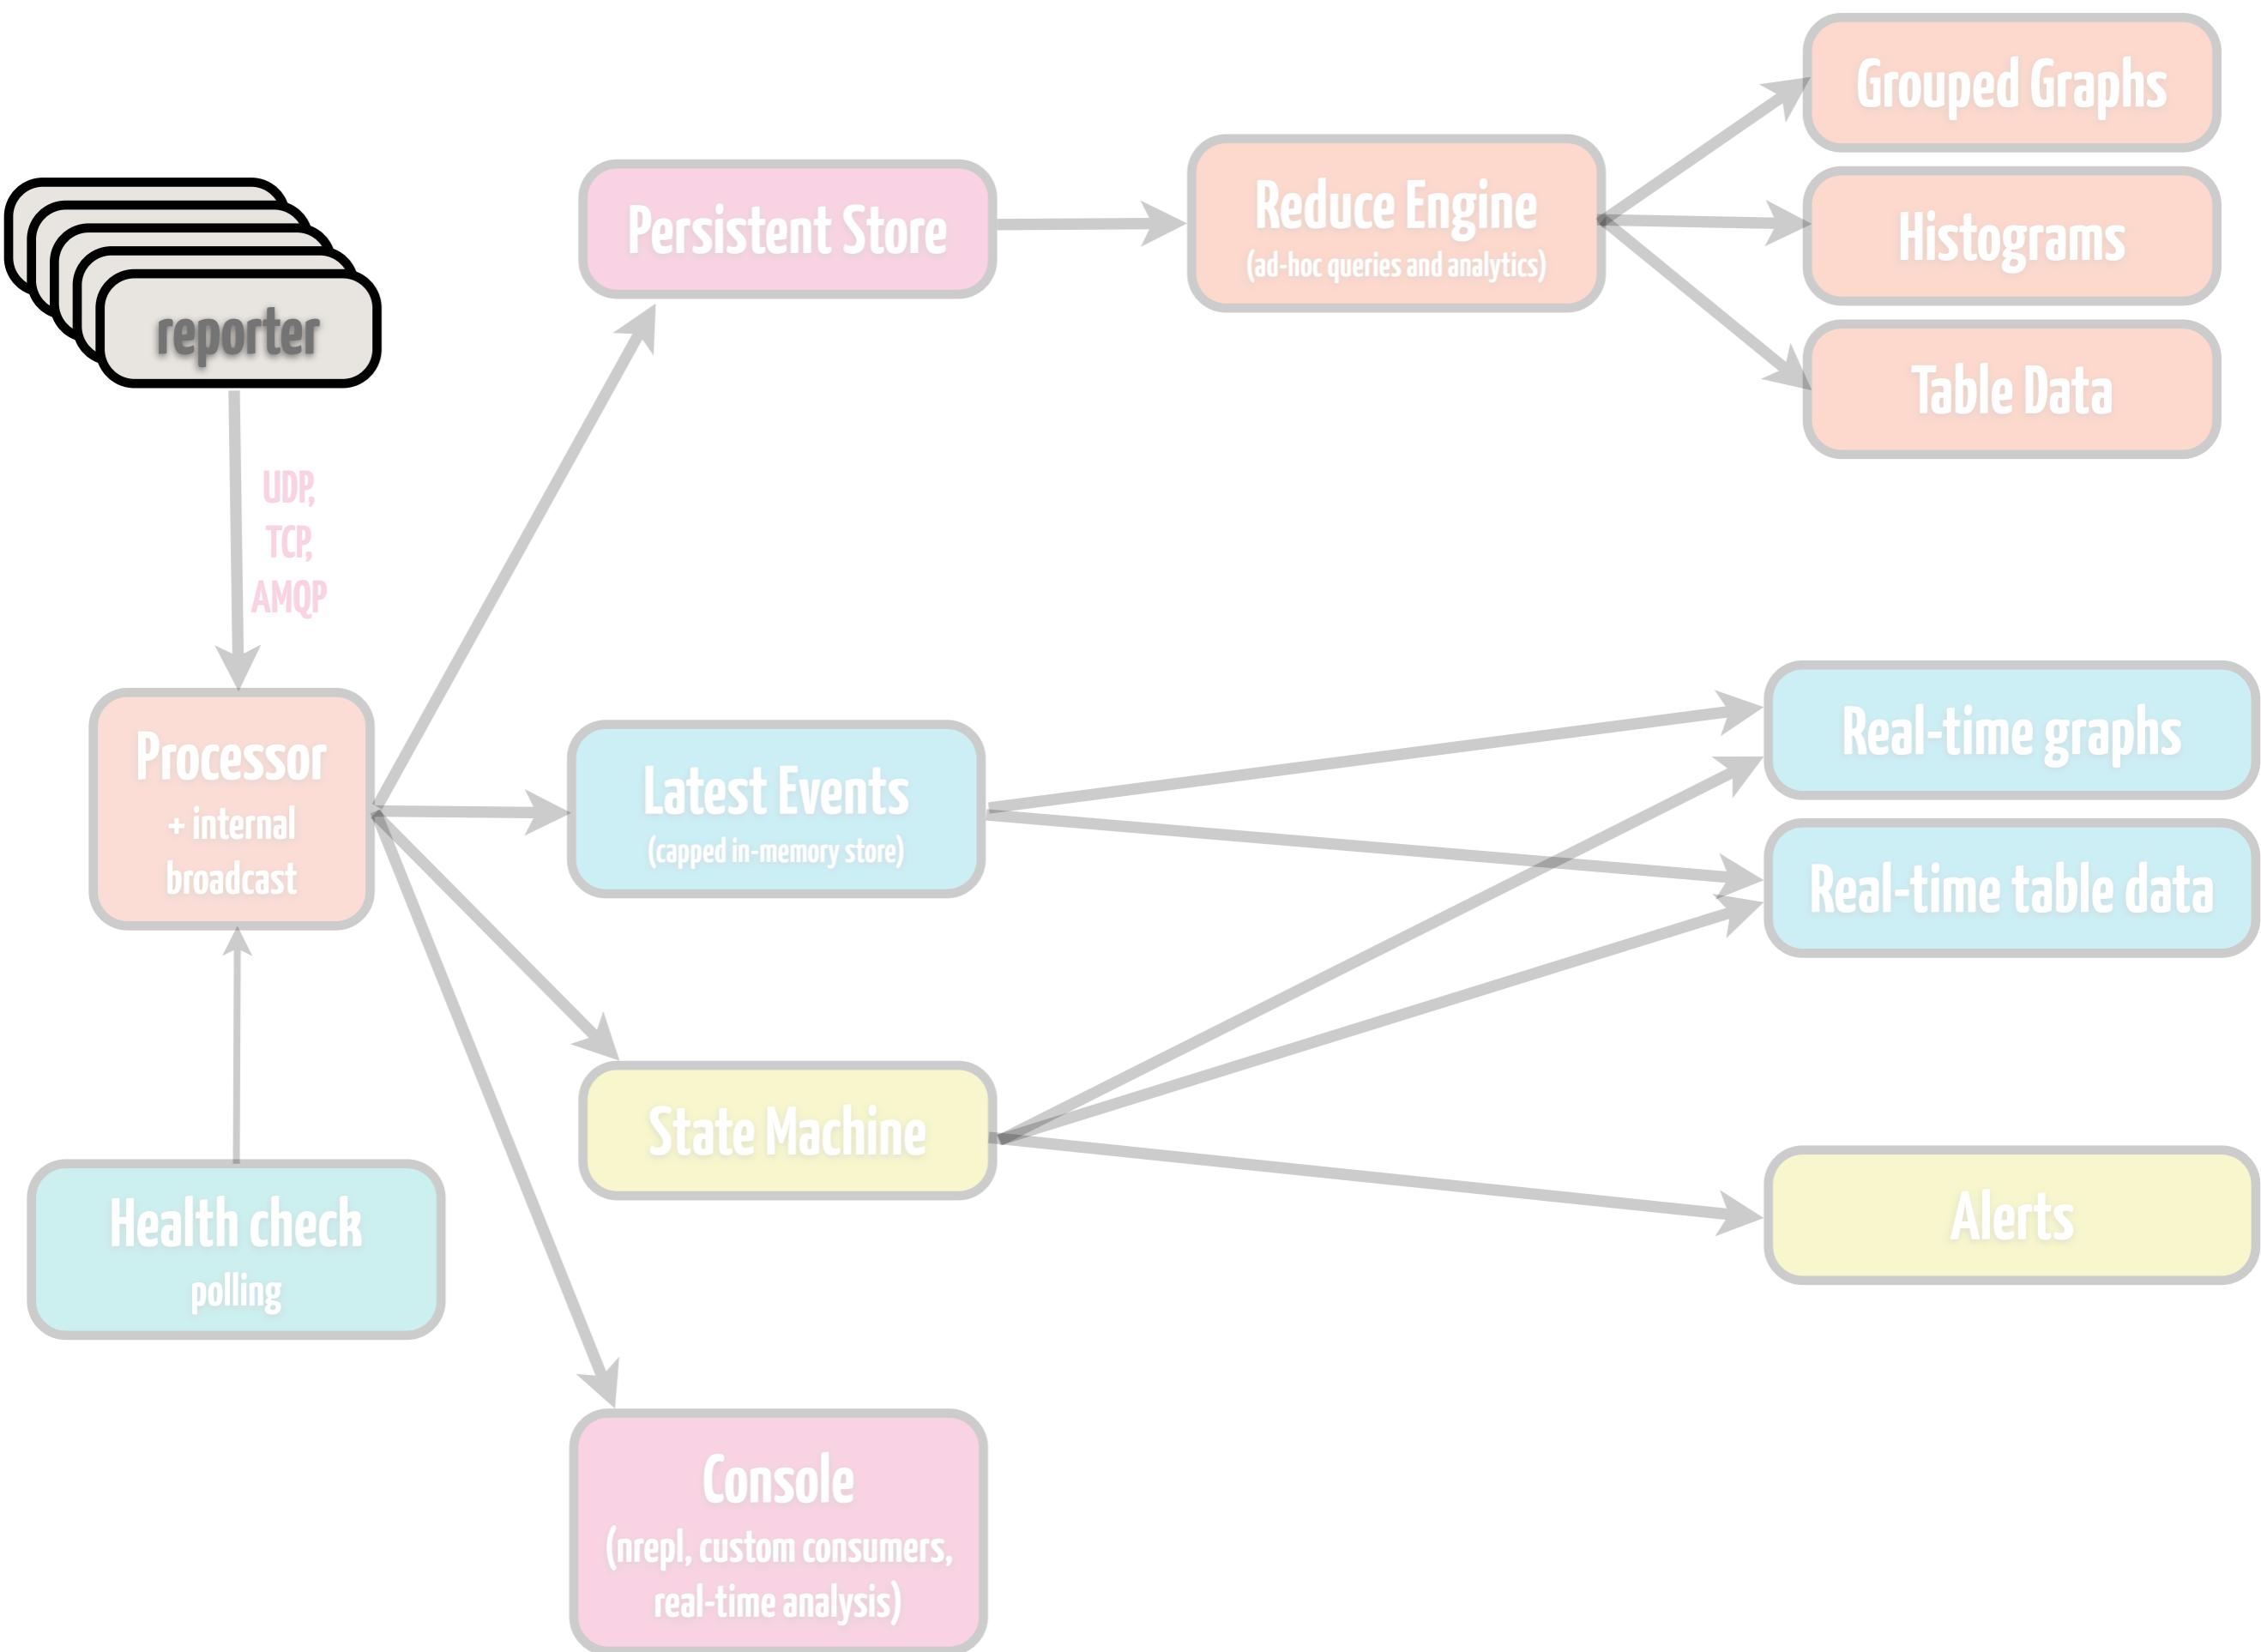
and try our best to make it **not**
opinionated

opinionated stuff is about tight
coupling

let's think of it in terms of a
toolchain, not a **framework** (tm)



parts



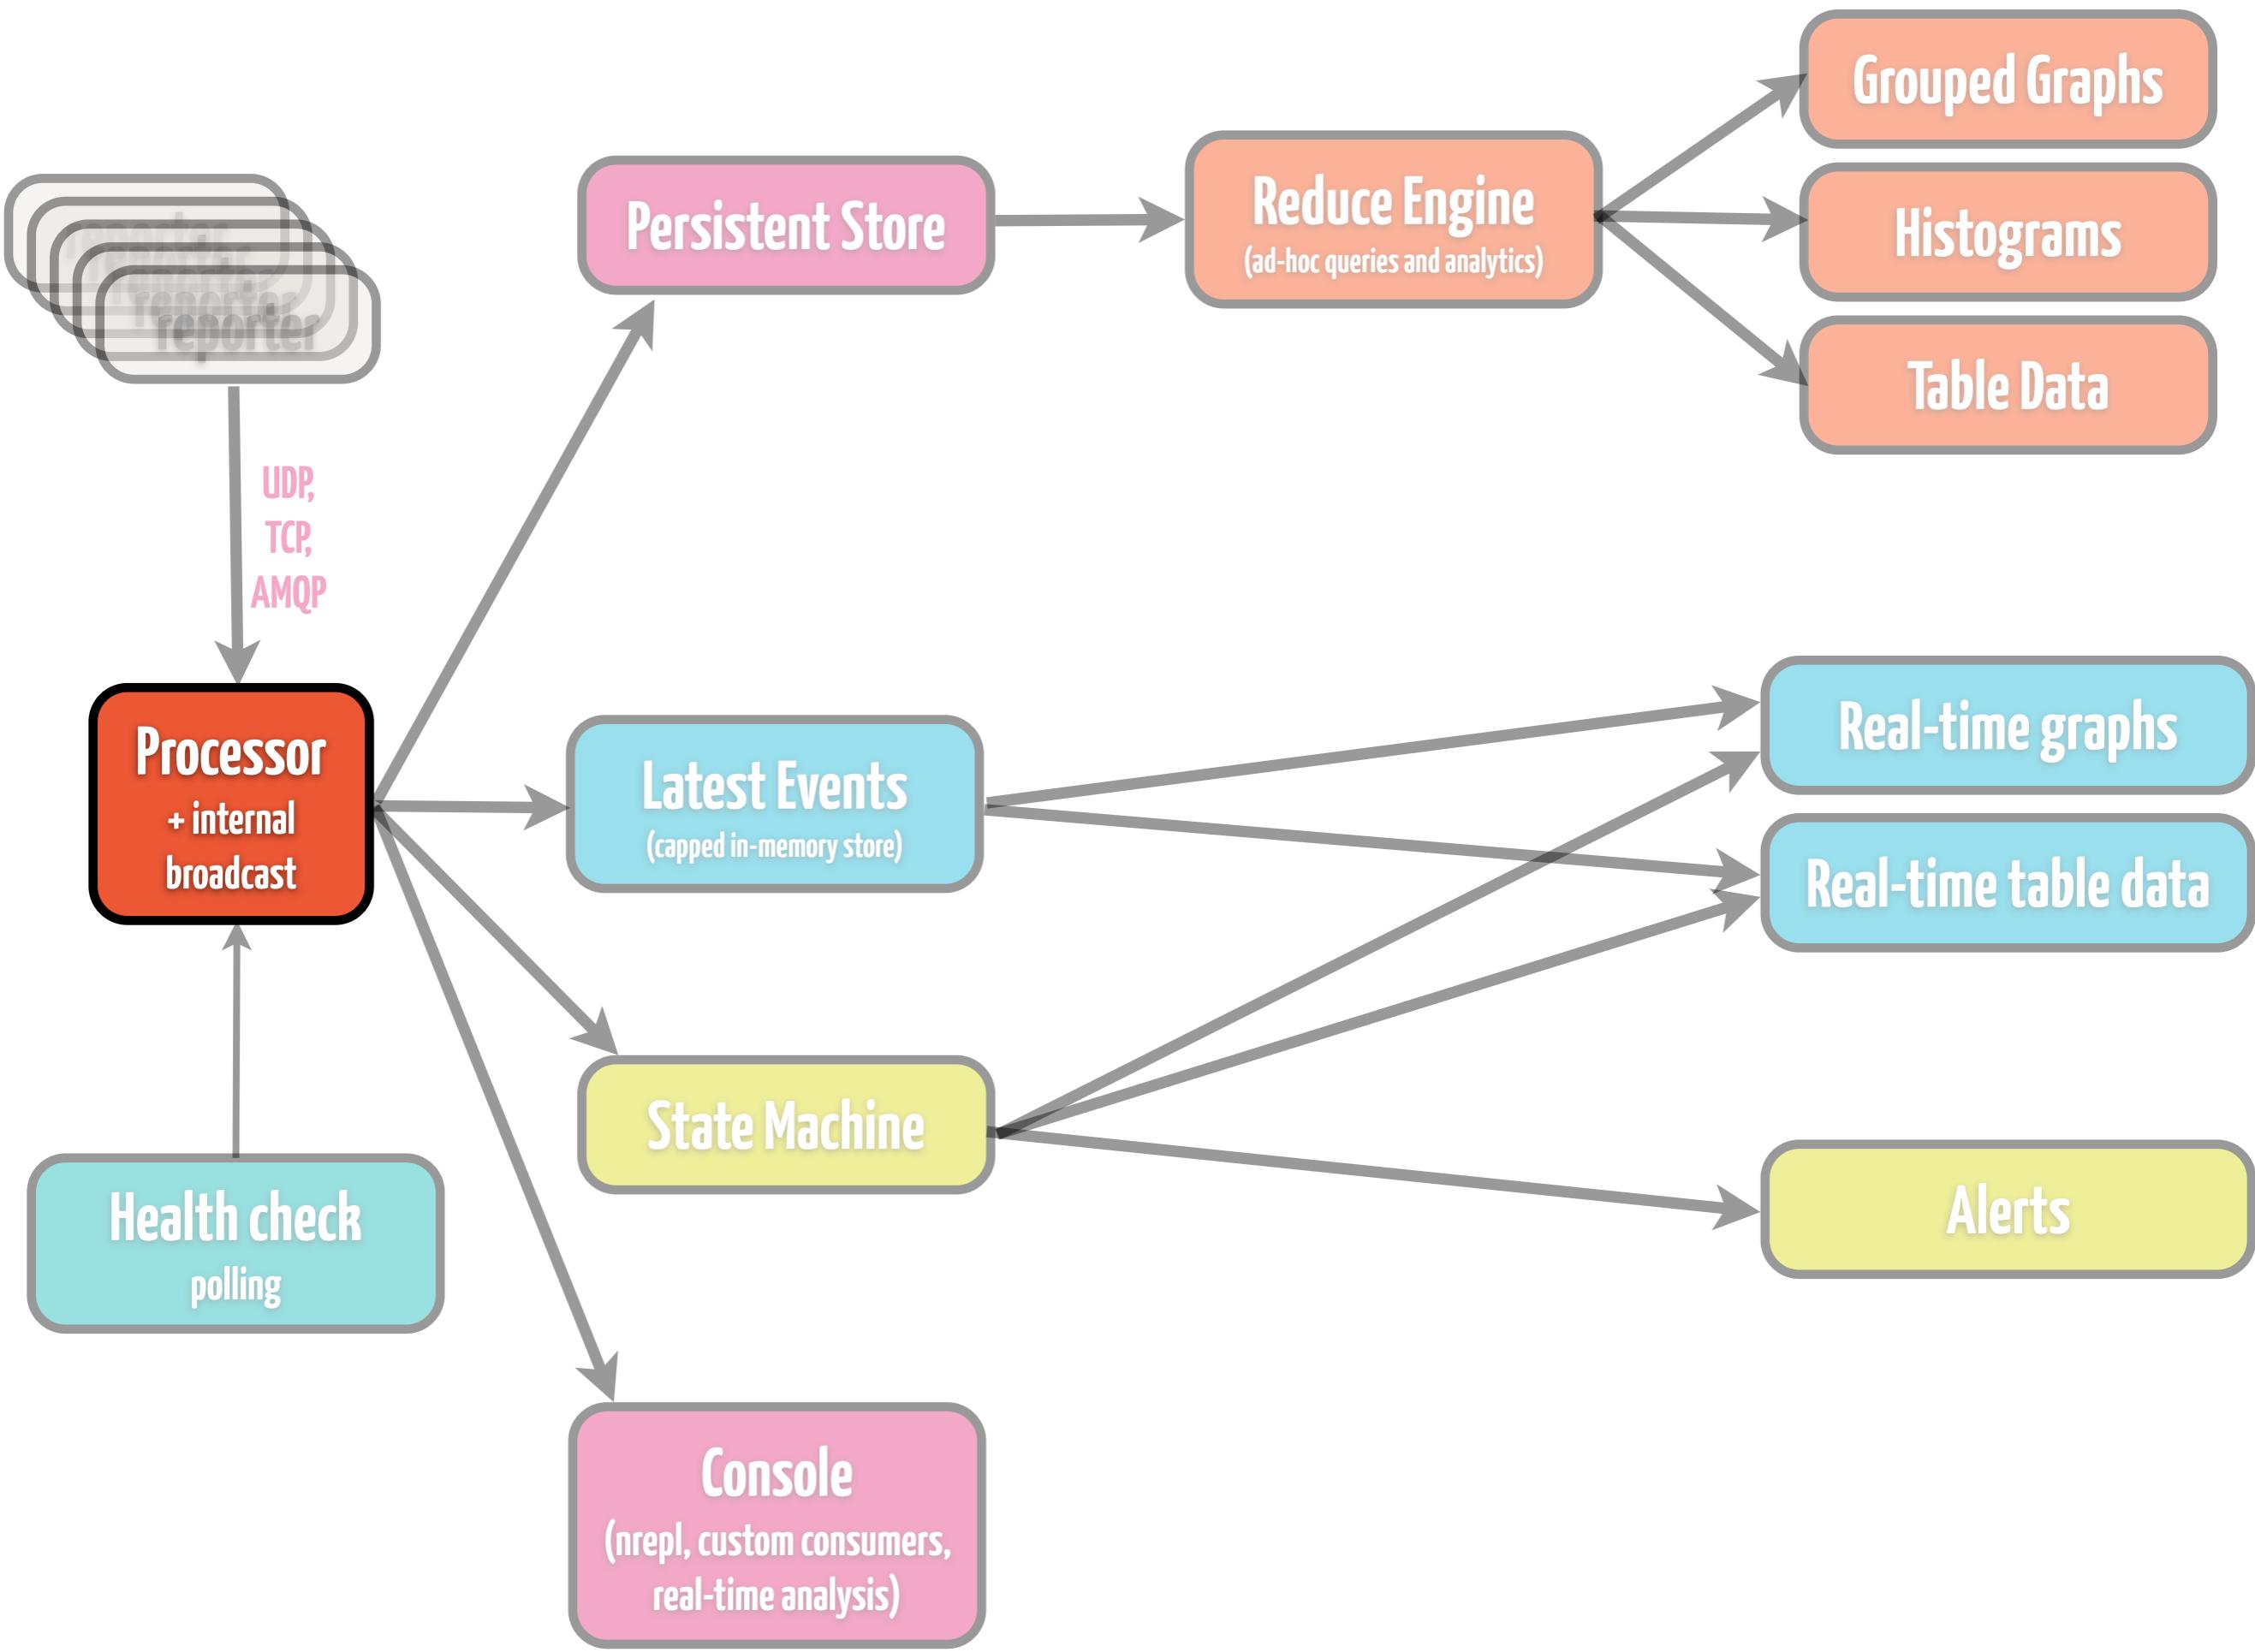


**simple, < 20 LOC software that
sends events to the server**



message format

```
{  
    ;; Md5 hash of an event, that uniquely identifies it  
    :md5 "cd0e351d2eefdf0f79e0b55a0efe543b",  
    ;; Arbitrary additional info  
    :additional_info  
        {  
            :execution_time 0.546999,  
            :url "http://mysite.com/page0" },  
    ;; Event Type identifier (404, exception, page load time)  
    :type "page_load",  
    ;; Dispatcher host name  
    :hostname "dc0-web01",  
    ;; Time when event was received  
    :received_at ...,  
    ;; Tags assigned to the event  
    :tags ["metrics" "performance"]  
}
```



Processor
+ internal
broadcast

In nutshell, it's a pipeline.

Basic properties:

- validate the event
- classify / re-classify event
- run processor based on current event type
- add calculated params
- distribute events internally

Processor
+ internal
broadcast

Classification

- remember the example with 404 that looked like an exception?
- makes events harvesting and further processing easier
- client will remain dumb, no need to redeploy on type changes
- useful for sub-typing, e.g. page_load could become slow_page_load etc.
- depends on your use case



message format

```
{  
  :additional_info  
    { :execution_time 0.5469999,  
      :url "http://mysite.com/page0" },  
  
  :hostname "dc0-web01",  
  :received_at ....,  
  :tags ["metrics" "performance"]  
}
```



message format

```
{  
  :additional_info  
    { :execution_time 0.5469999,  
      :url "http://mysite.com/page0" },  
  :type "page_load",  
  :hostname "dc0-web01",  
  :received_at ....,  
  :tags ["metrics" "performance"]  
}
```



message format

```
{  
  :additional_info  
    { :execution_time 0.546999,  
      :url "http://mysite.com/" },  
  
  :hostname "dc0-web01",  
  :received_at ....,  
  :tags ["metrics" "performance"]  
}
```



message format

```
{  
  :additional_info  
    { :execution_time 0.5469999,  
      :url "http://mysite.com/" },  
  :type "landing_page_load",  
  :hostname "dc0-web01",  
  :received_at ....,  
  :tags ["metrics" "performance"]  
}
```

Processor
+ internal
broadcast

Calculated params

- internal example is md5 hash, used to ease finding unique events
- could be used for grouping within a type
- percent of a whole calculation max heap: 4gb, current: 2gb => 50% of a whole
- add flags for further pipeline modules



message format

```
{  
  :additional_info  
    { :backtrace "..."},  
  :type "backtrace",  
  :hostname "dc0-web01",  
  :received_at ...,  
  :tags ["metrics" "performance"]  
}
```



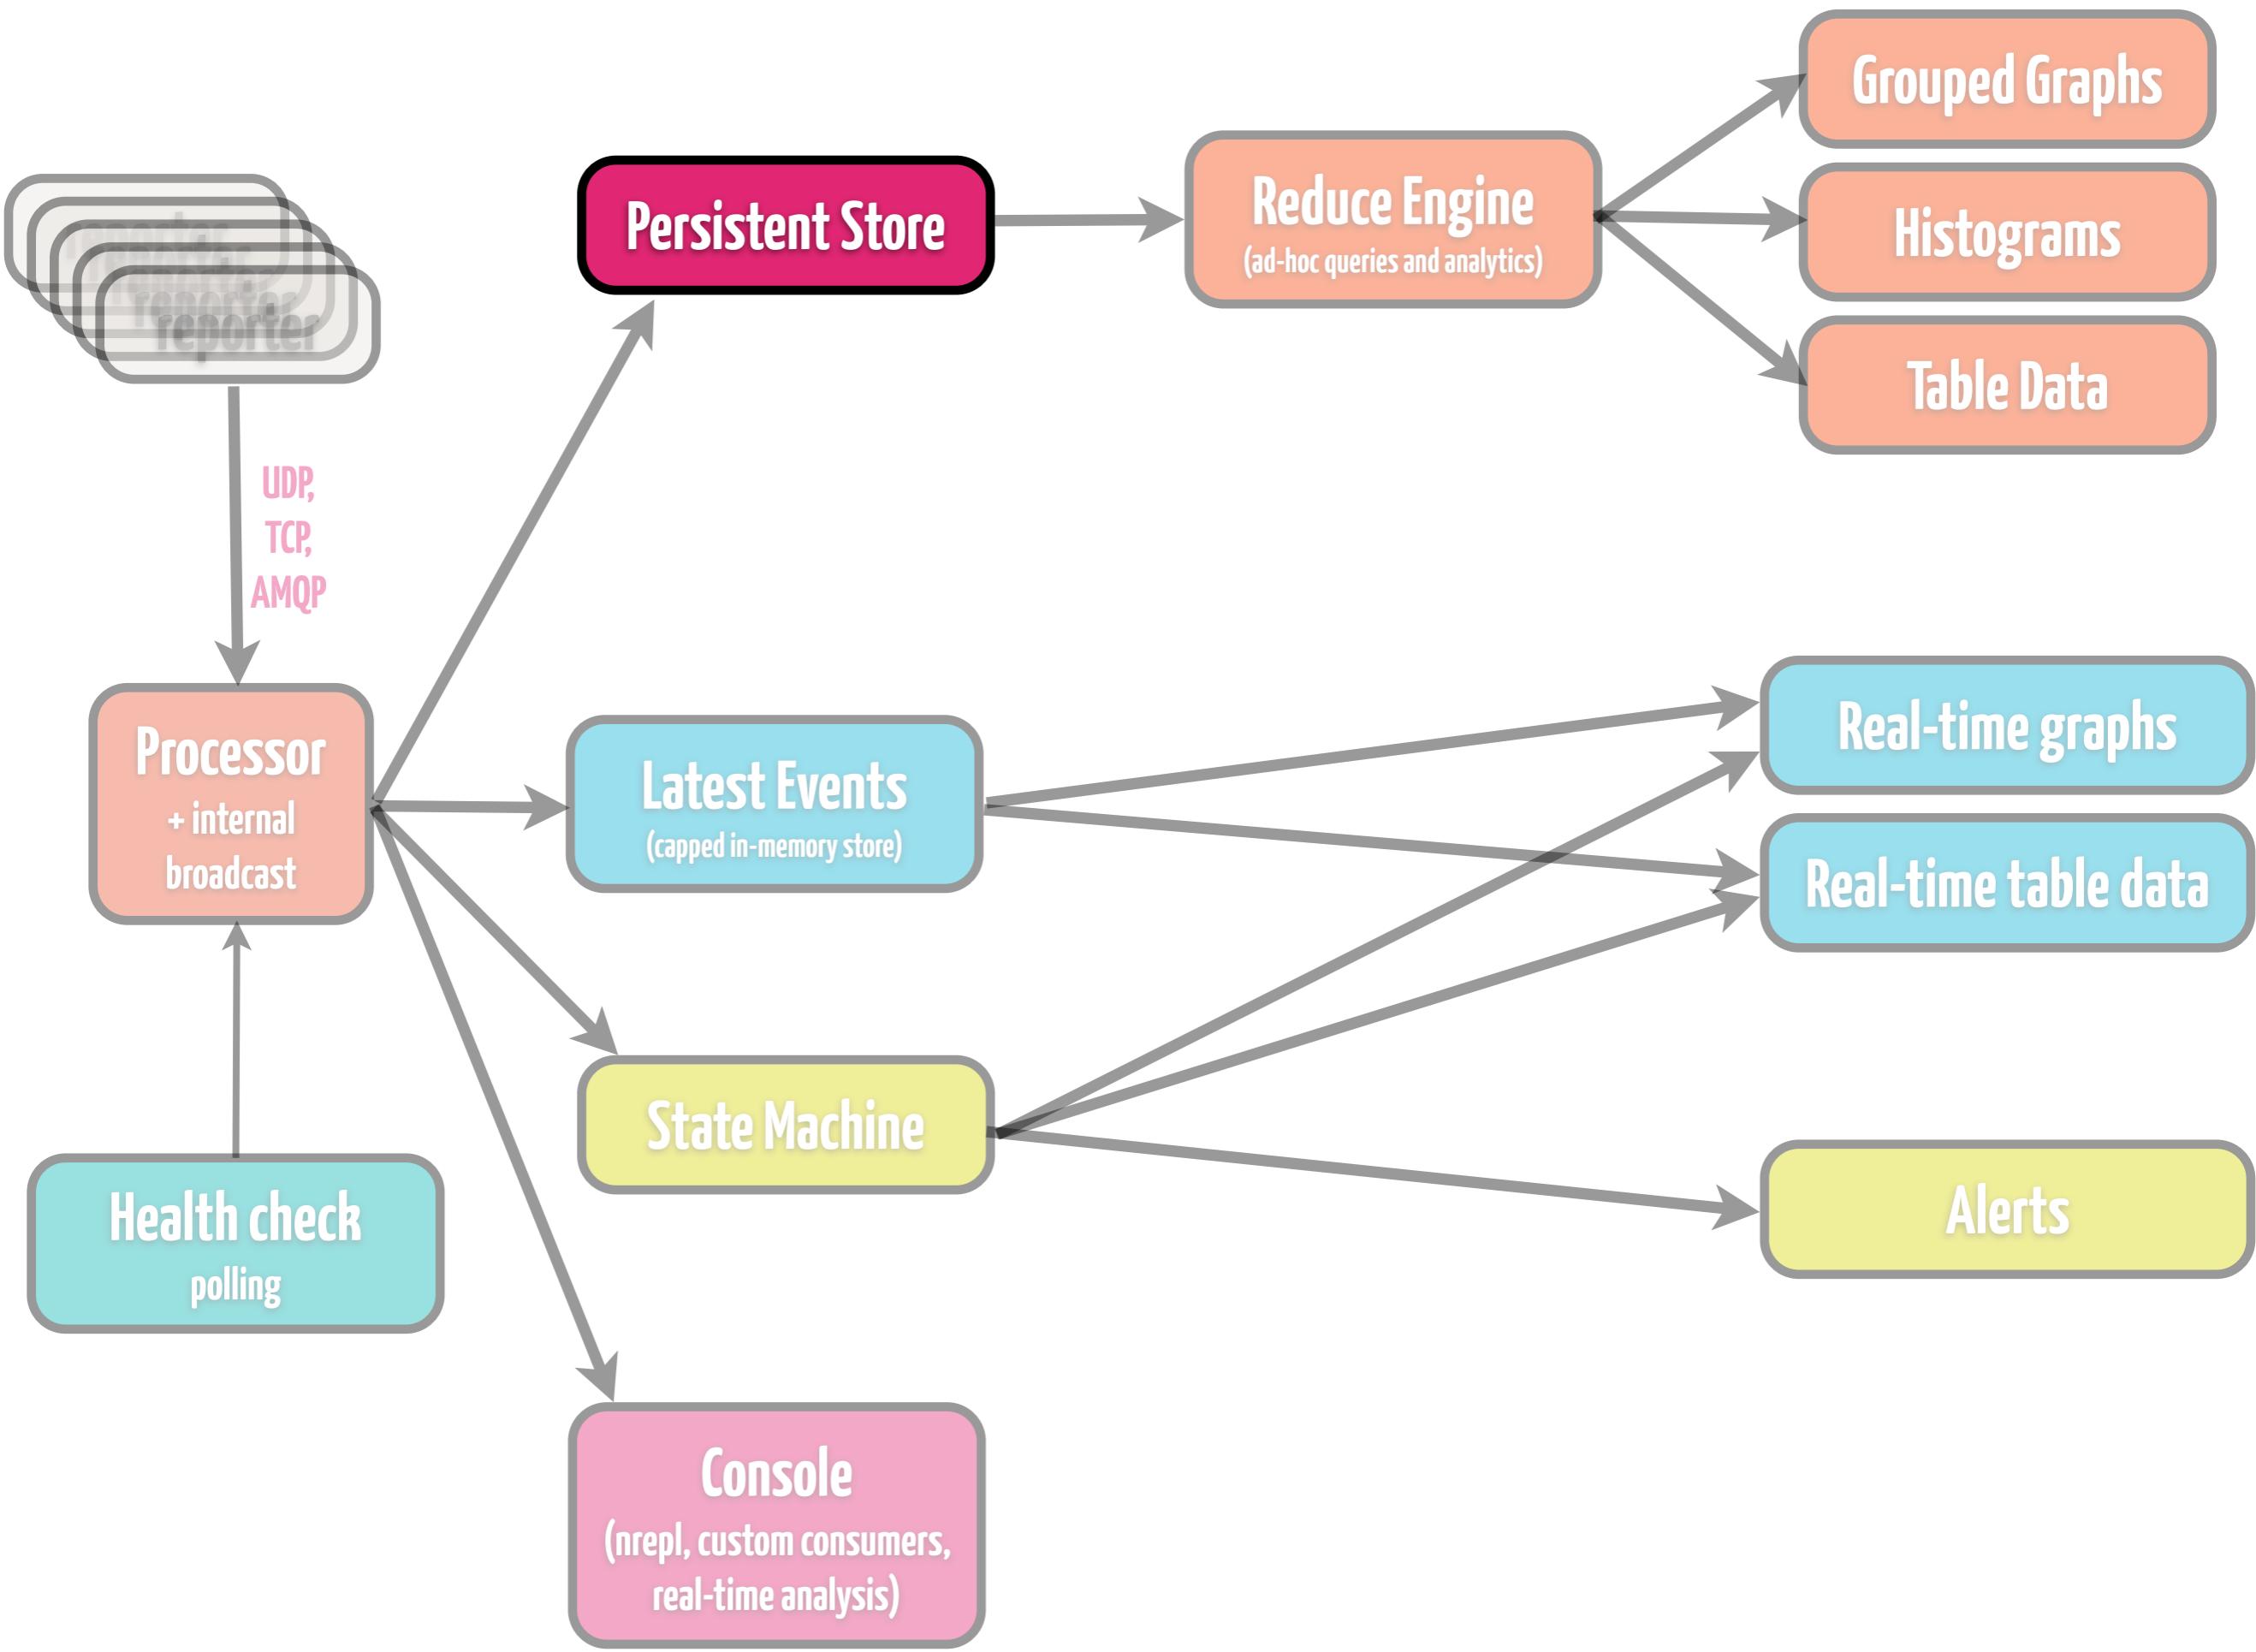
message format

```
{  
  :additional_info  
    { :backtrace "..."},  
  :md5 (calculate-md5 (:backtrace "..."))  
  :type "backtrace",  
  :hostname "dc0-web01",  
  :received_at ...,  
  :tags ["metrics" "performance"]  
}
```

Processor
+ internal
broadcast

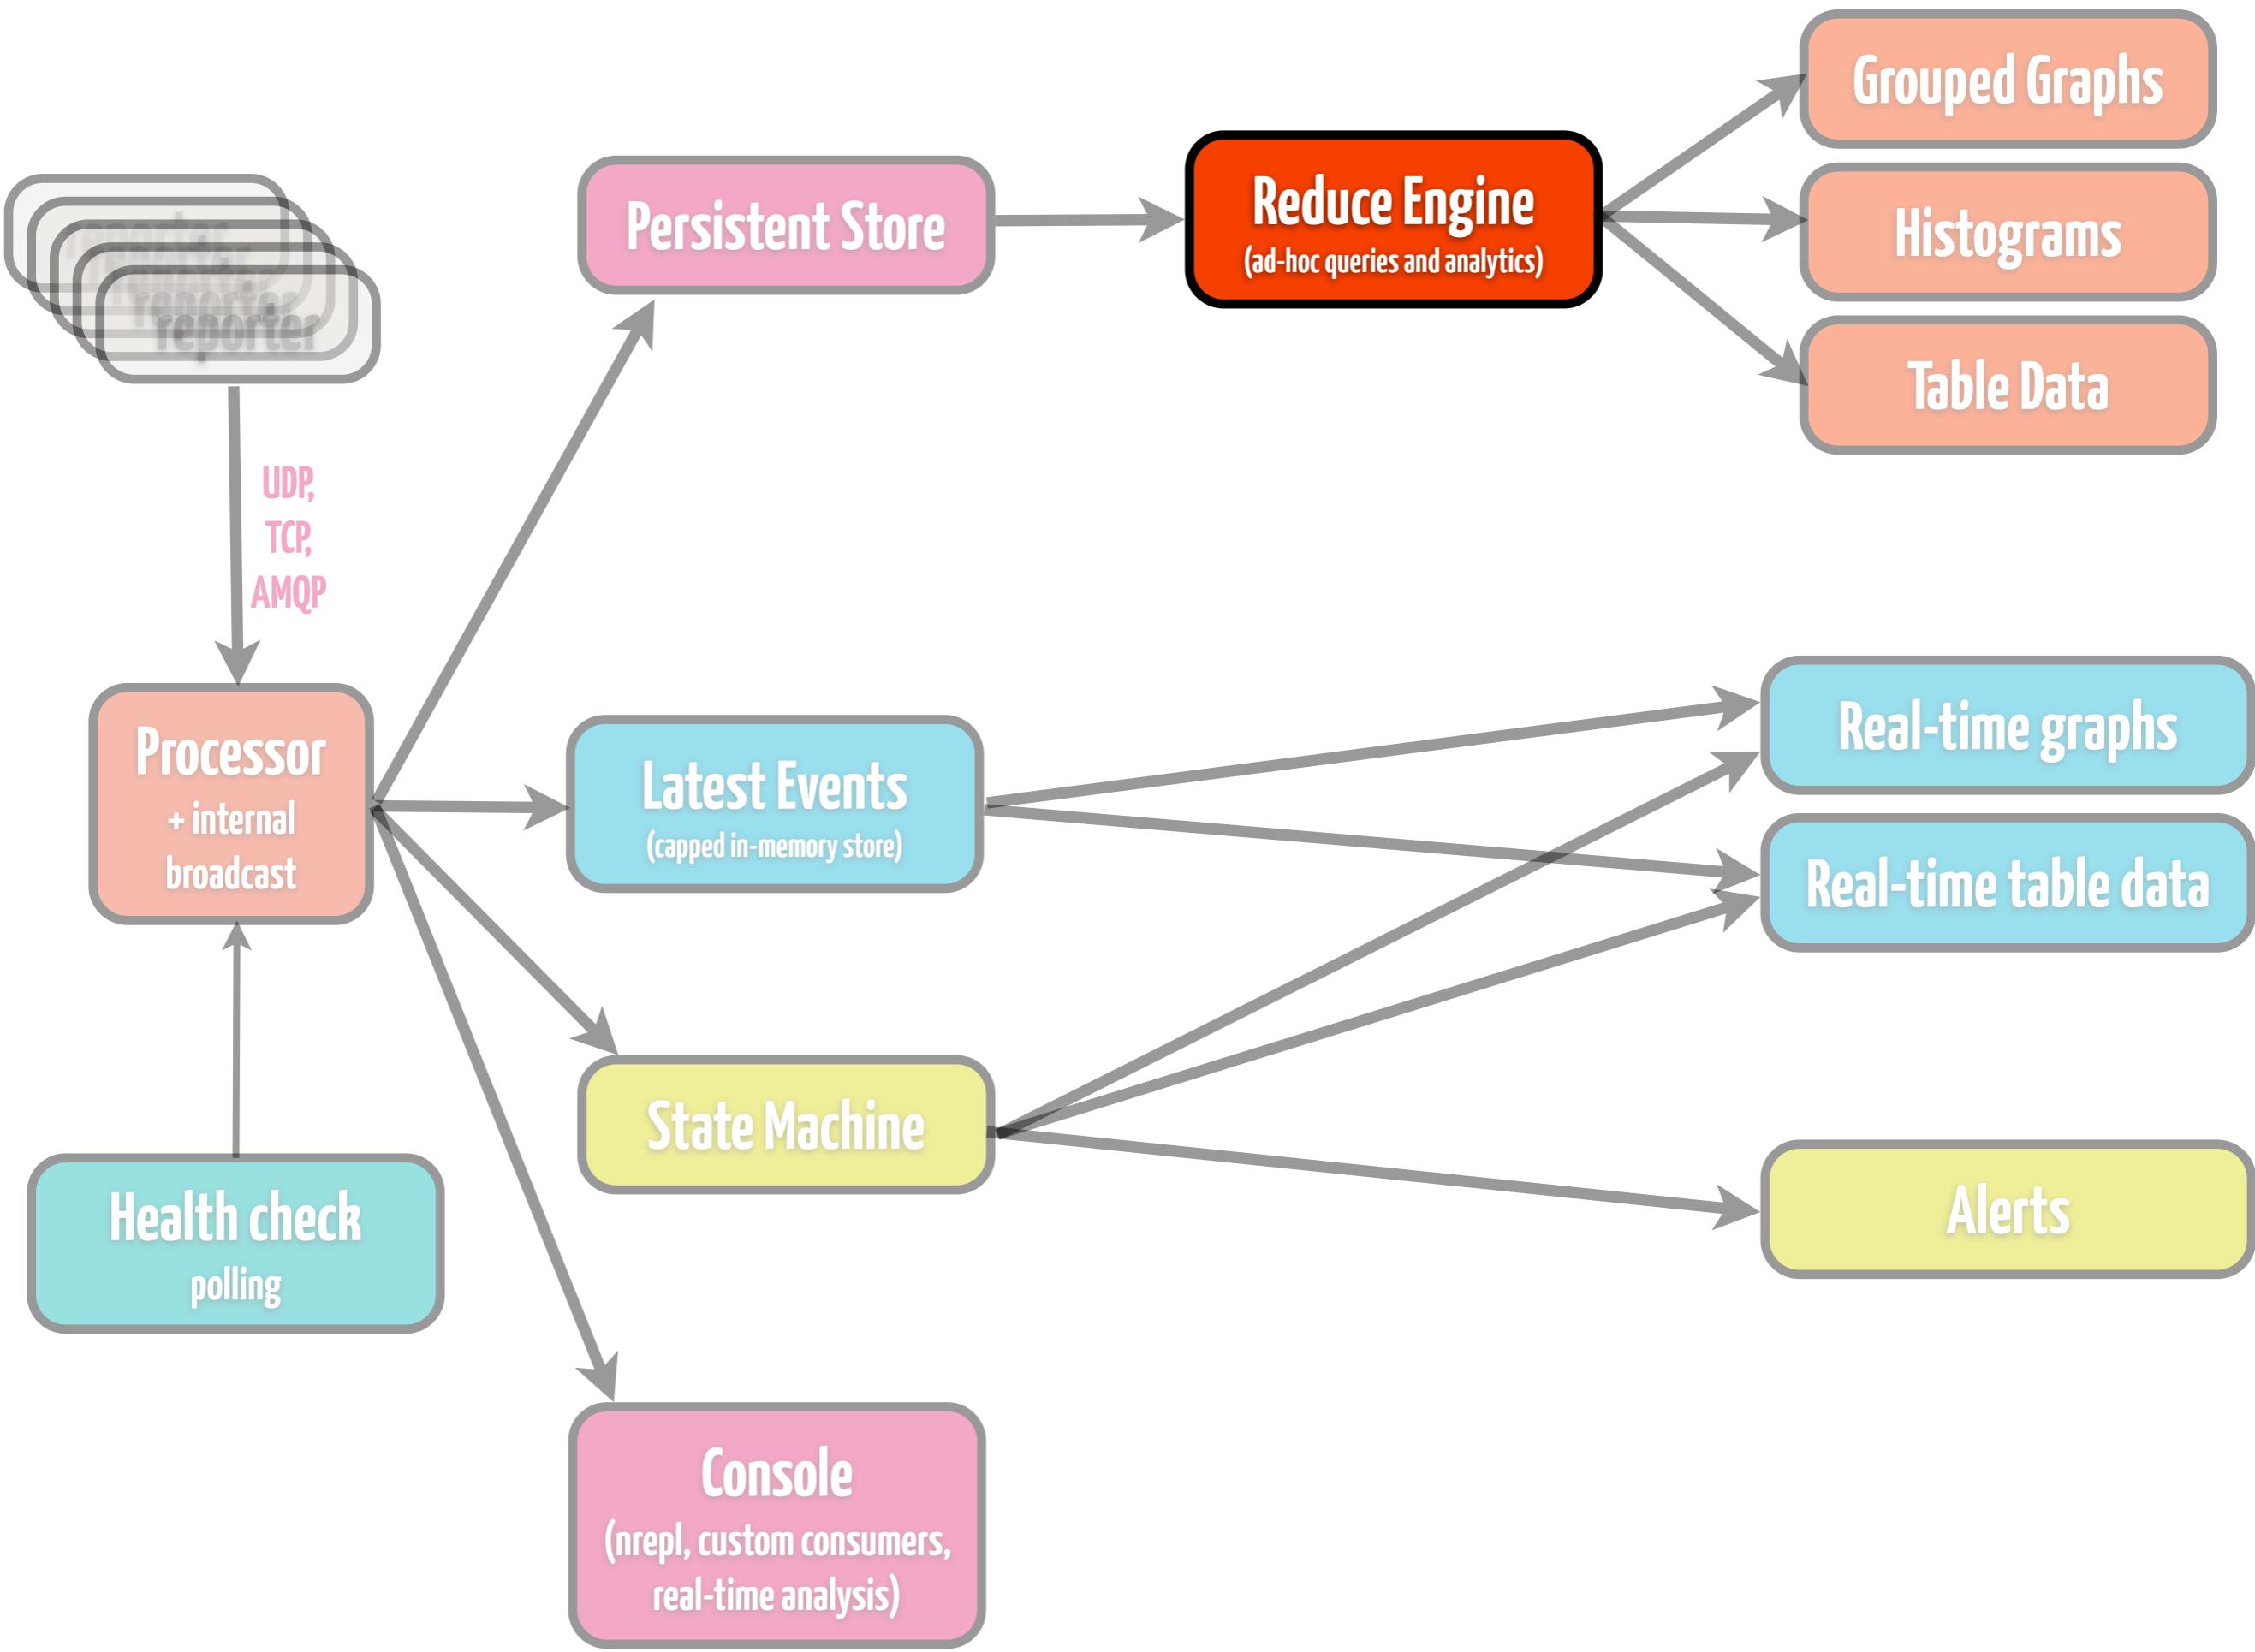
Internal broadcast

- when production system is running, you may want to attach to it in real-time and calculate message rate based on certain rules
- throw events to persistence layer, state machine, arbitrary listeners



Persistent Store

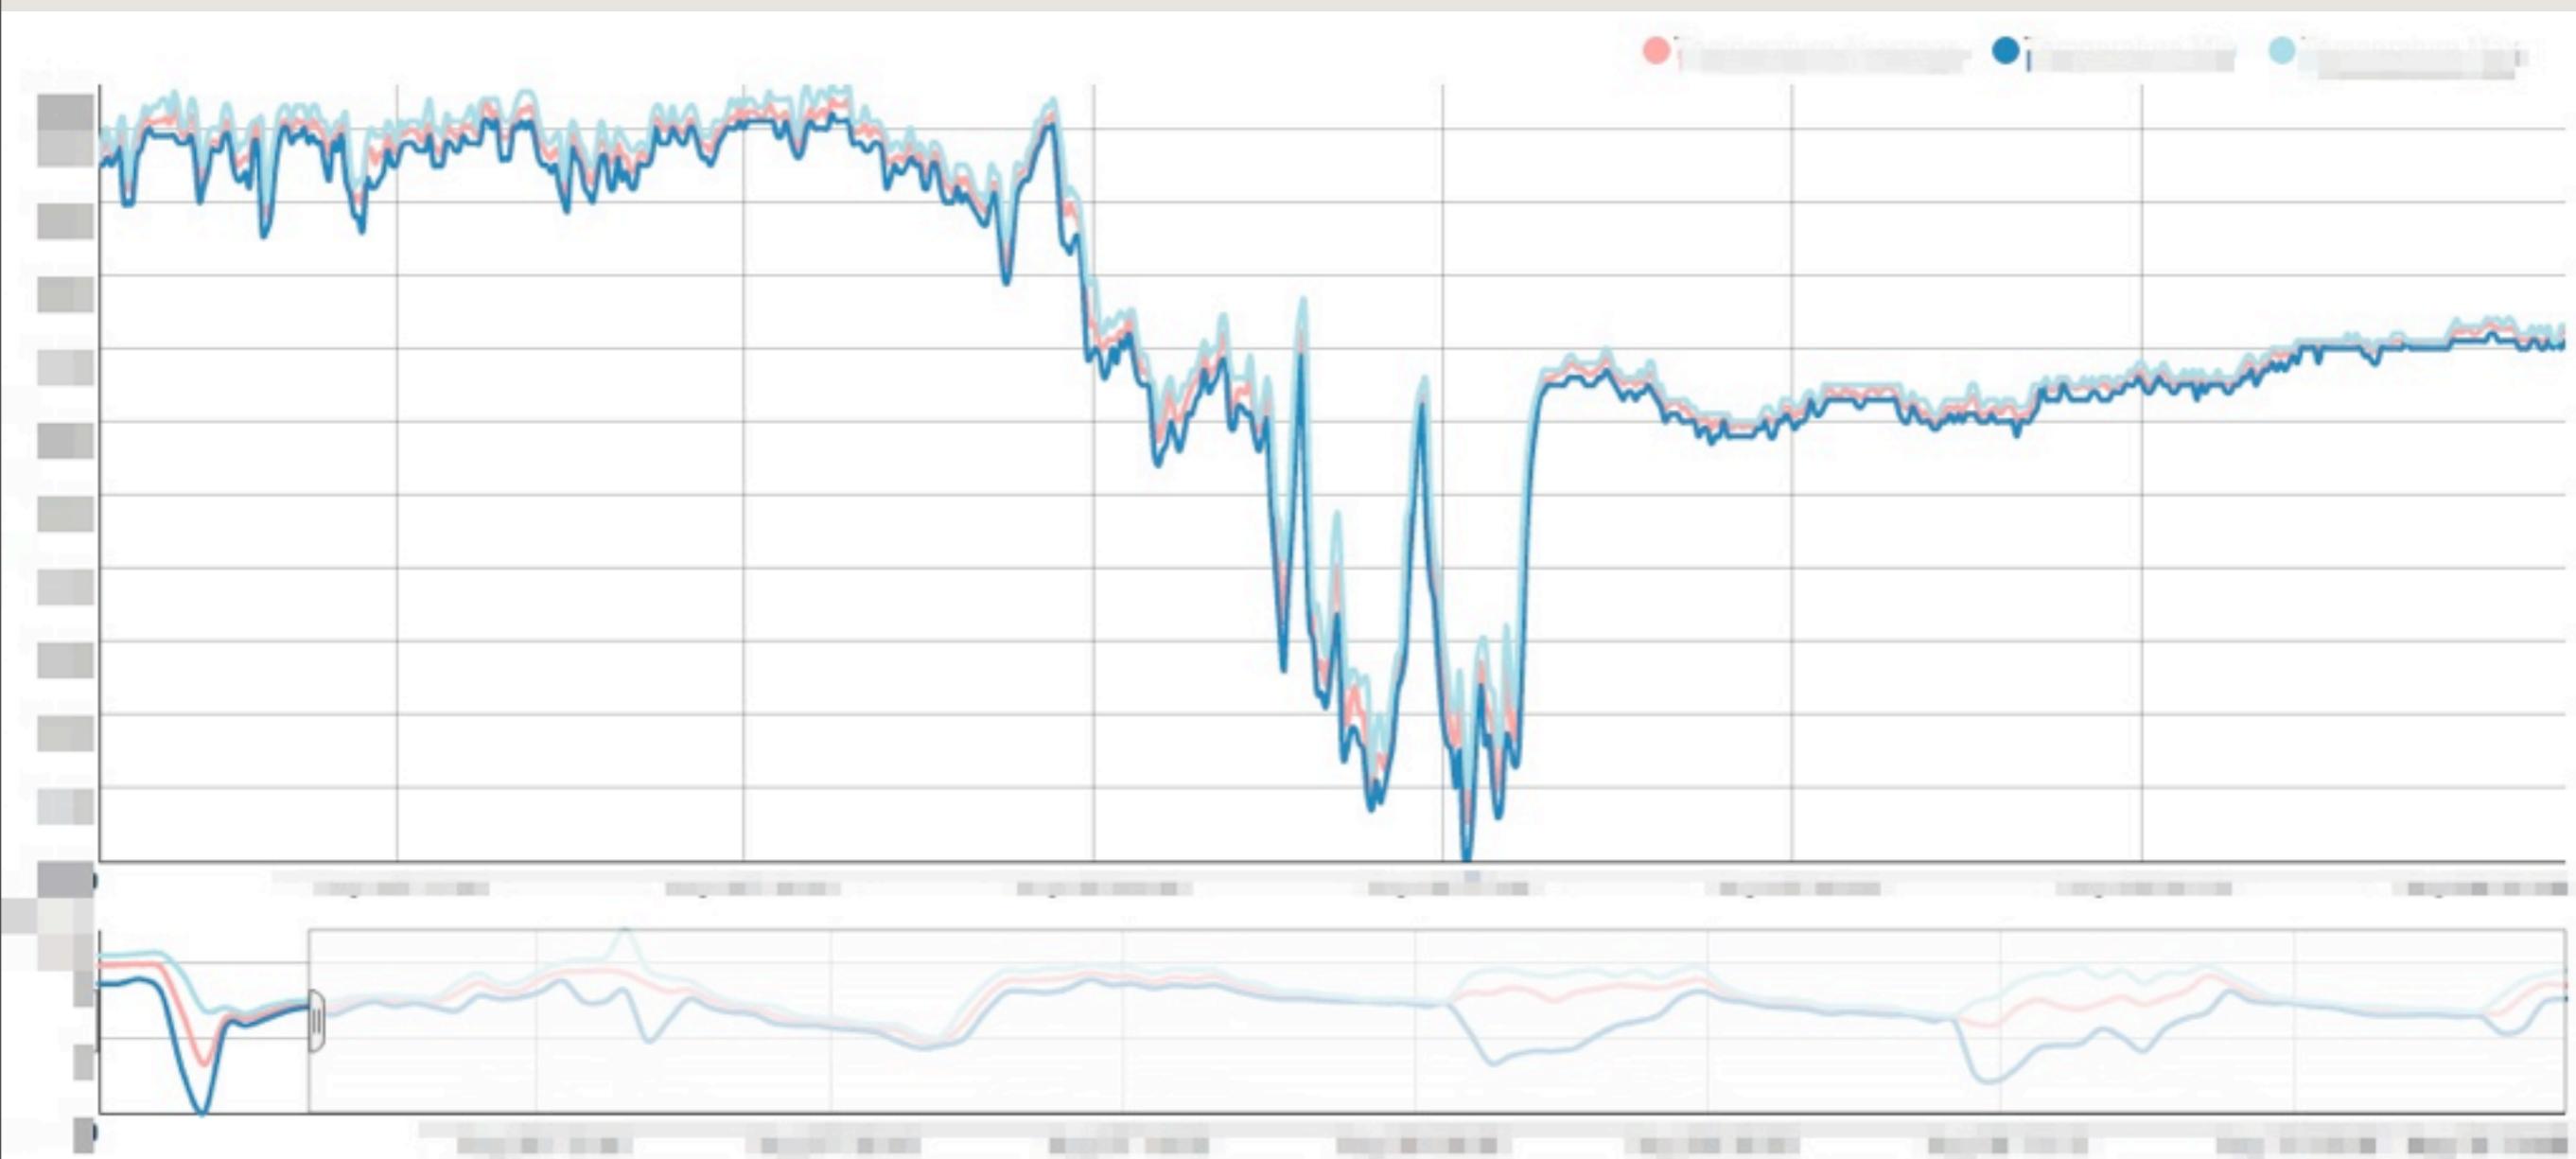
- highly depends on your throughput
- consider something lightweight and scalable
- if you generate fairly large amount of data, don't try to handle processing it by means of your store
- keep it simple
- think about Cassandra or HDFS, depending on how you plan to access data



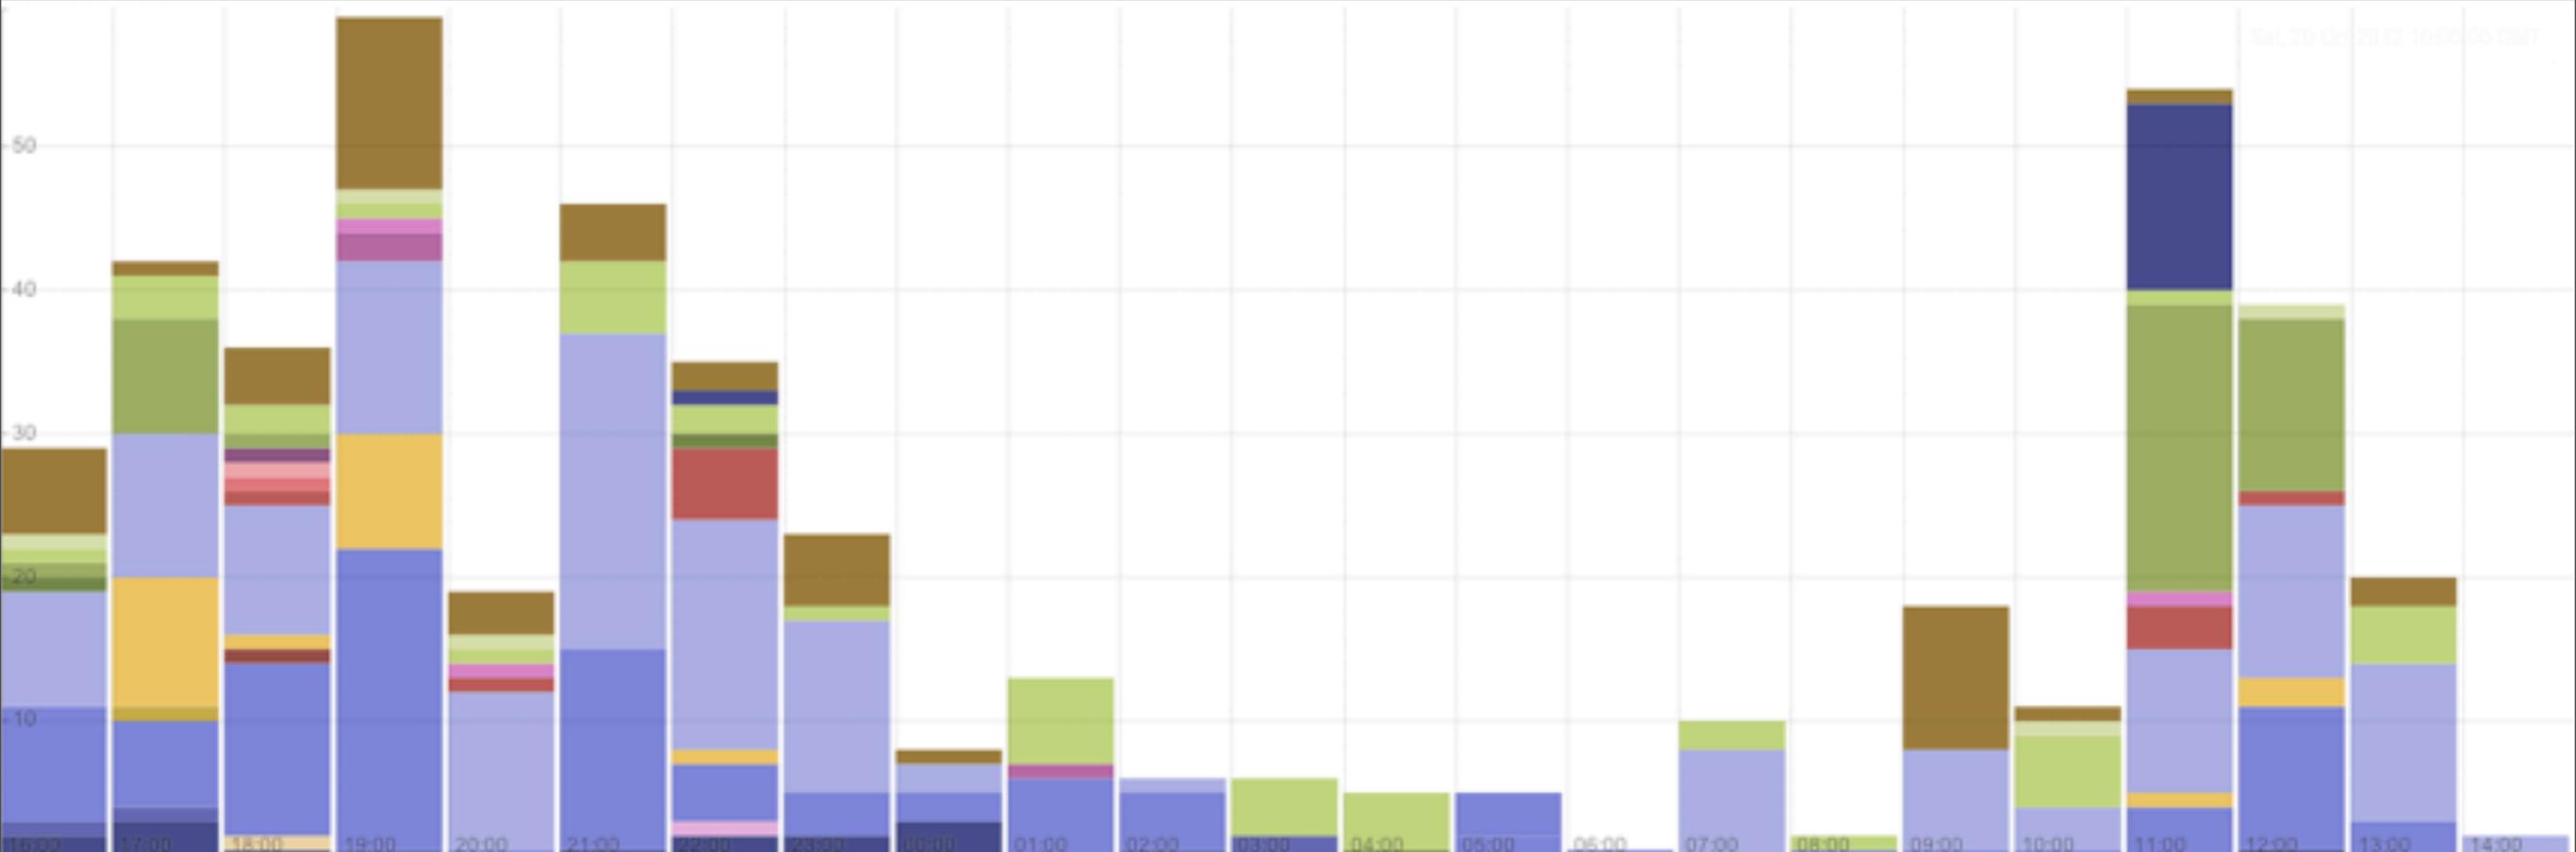
Reduce Engine
(ad-hoc queries and analytics)

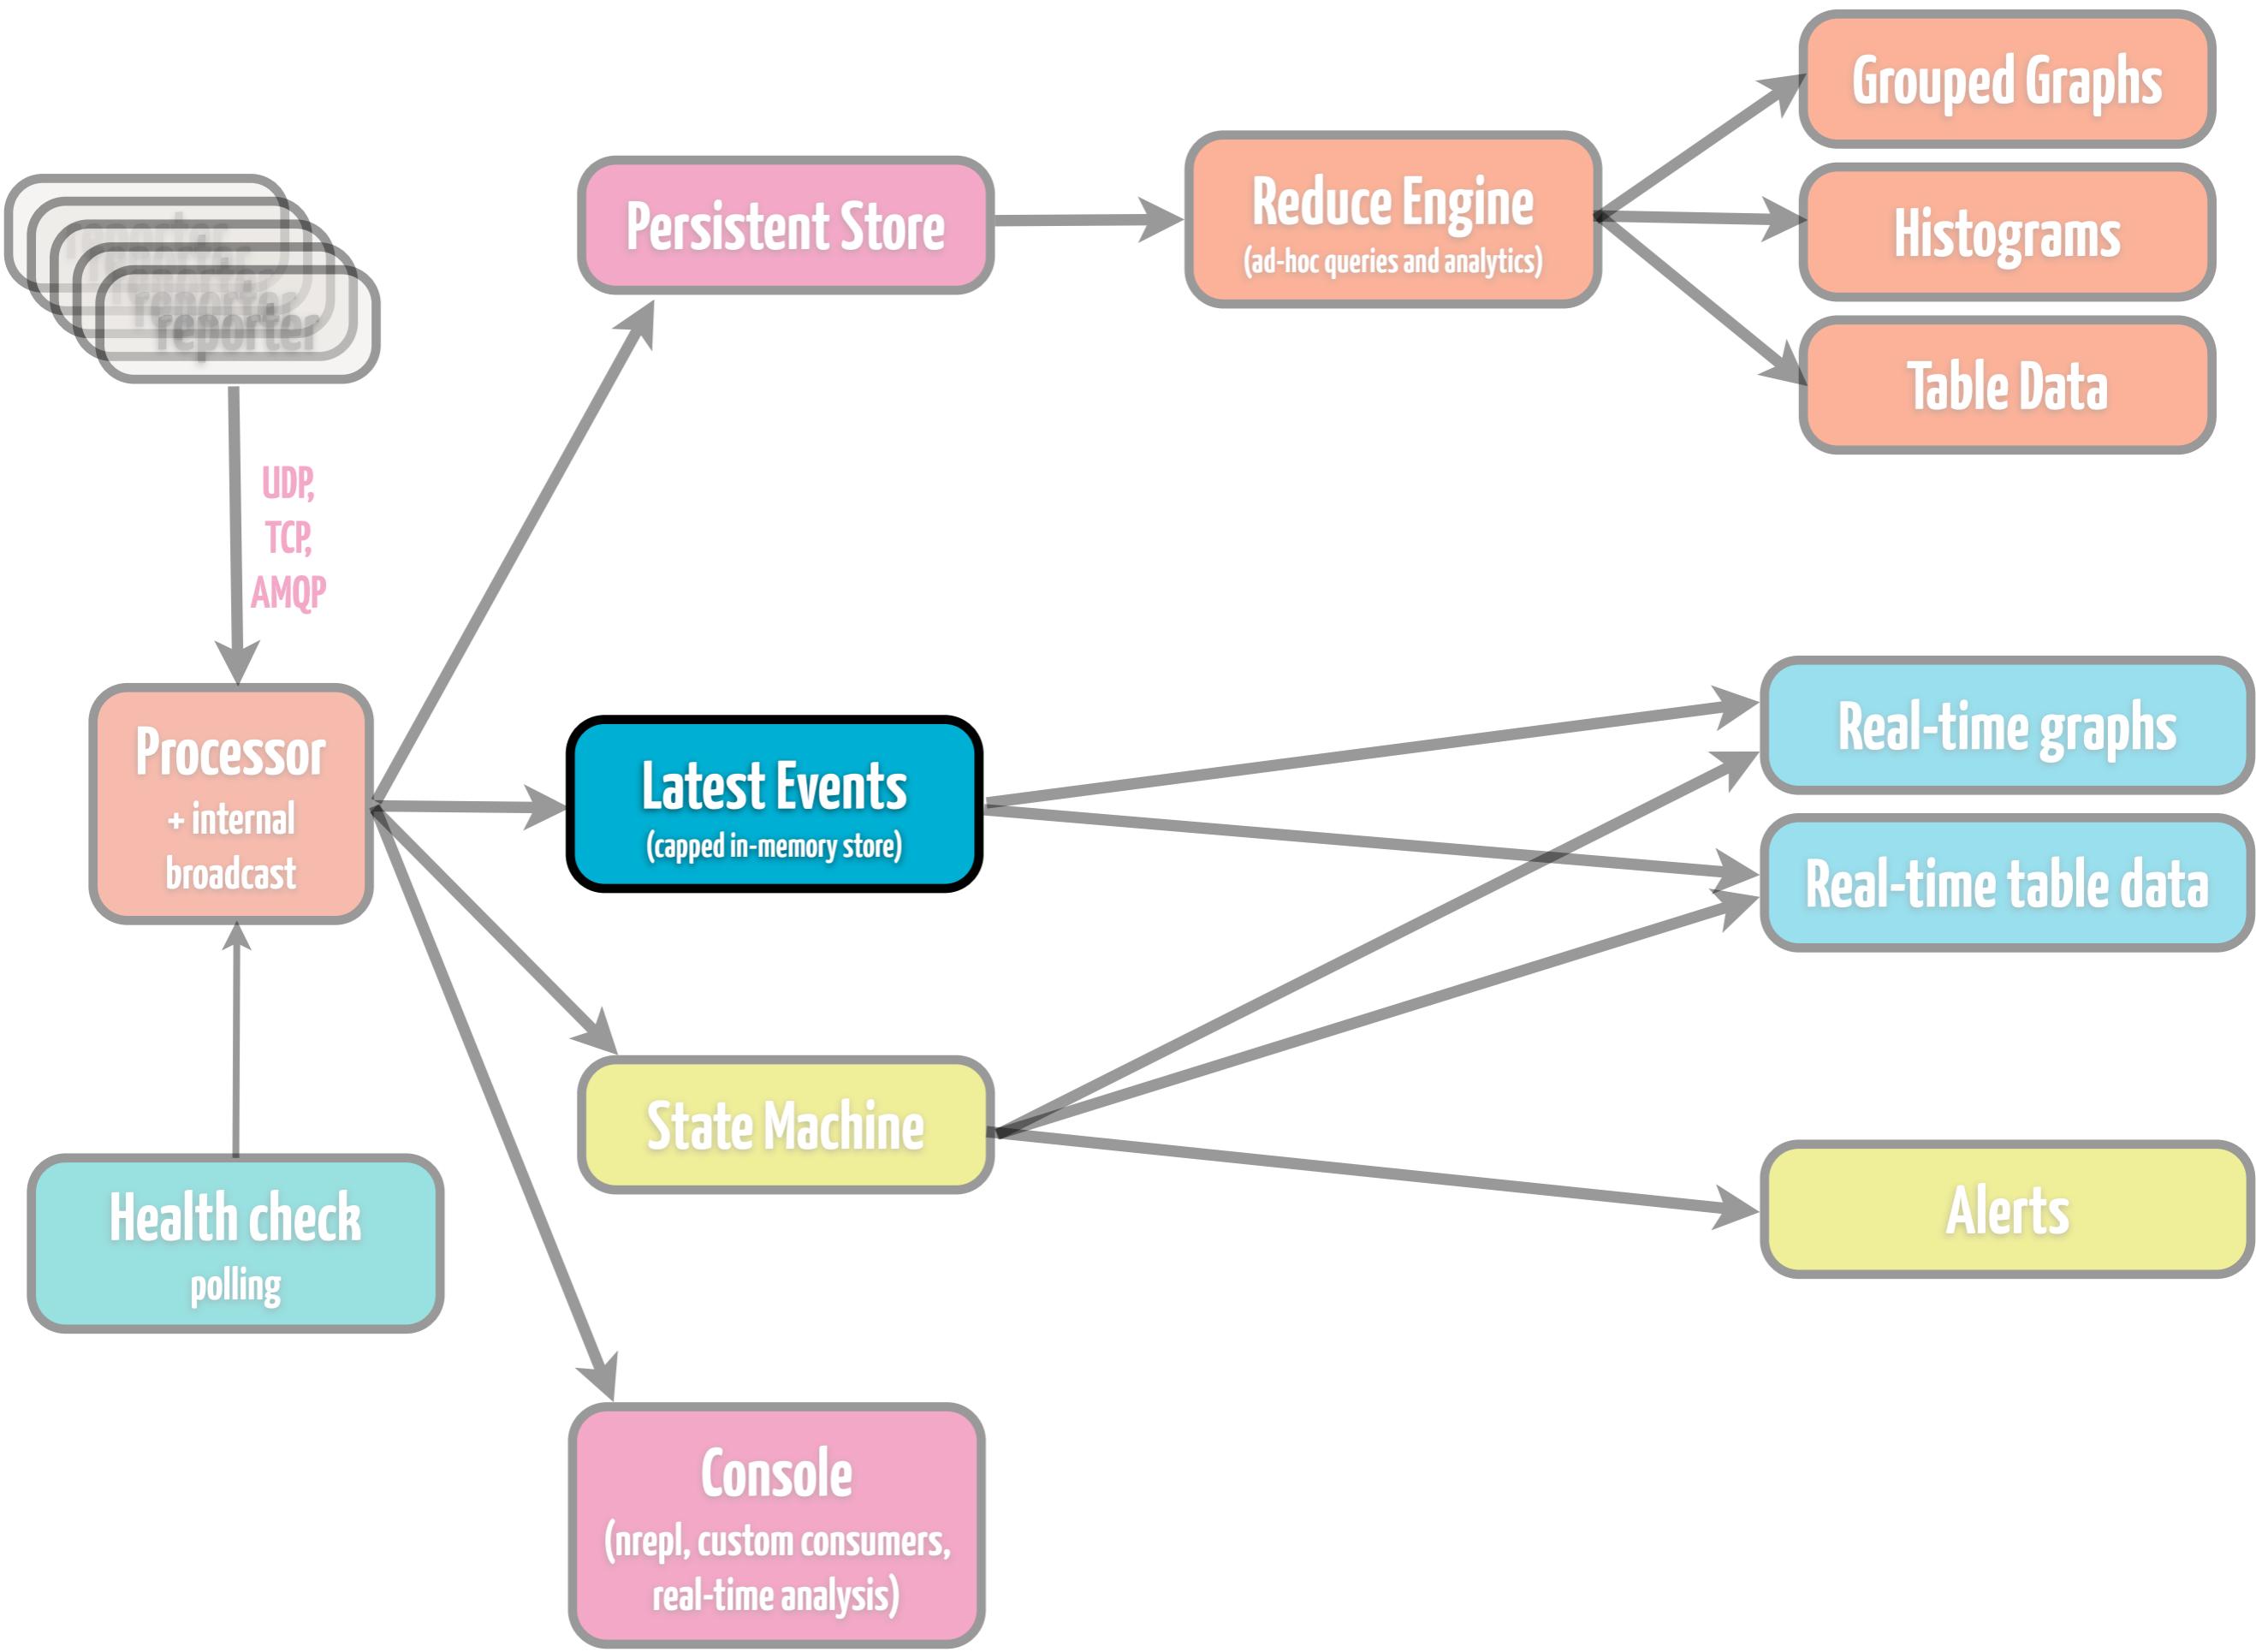
- sometimes you just don't know in real-time
- a new problem pops up, and you have to re-play stuff from a year ago
- find trends, that are simply not visible on a day/week/month scale
- most likely going to be something hadoop-based (YMMV)
- cascalog works quite well for us
 - ad-hoc queries
 - scheduled analytics
 - correlating metrics (hard to do in real time)

Graphs



Graphs





Latest Events
(capped in-memory store)

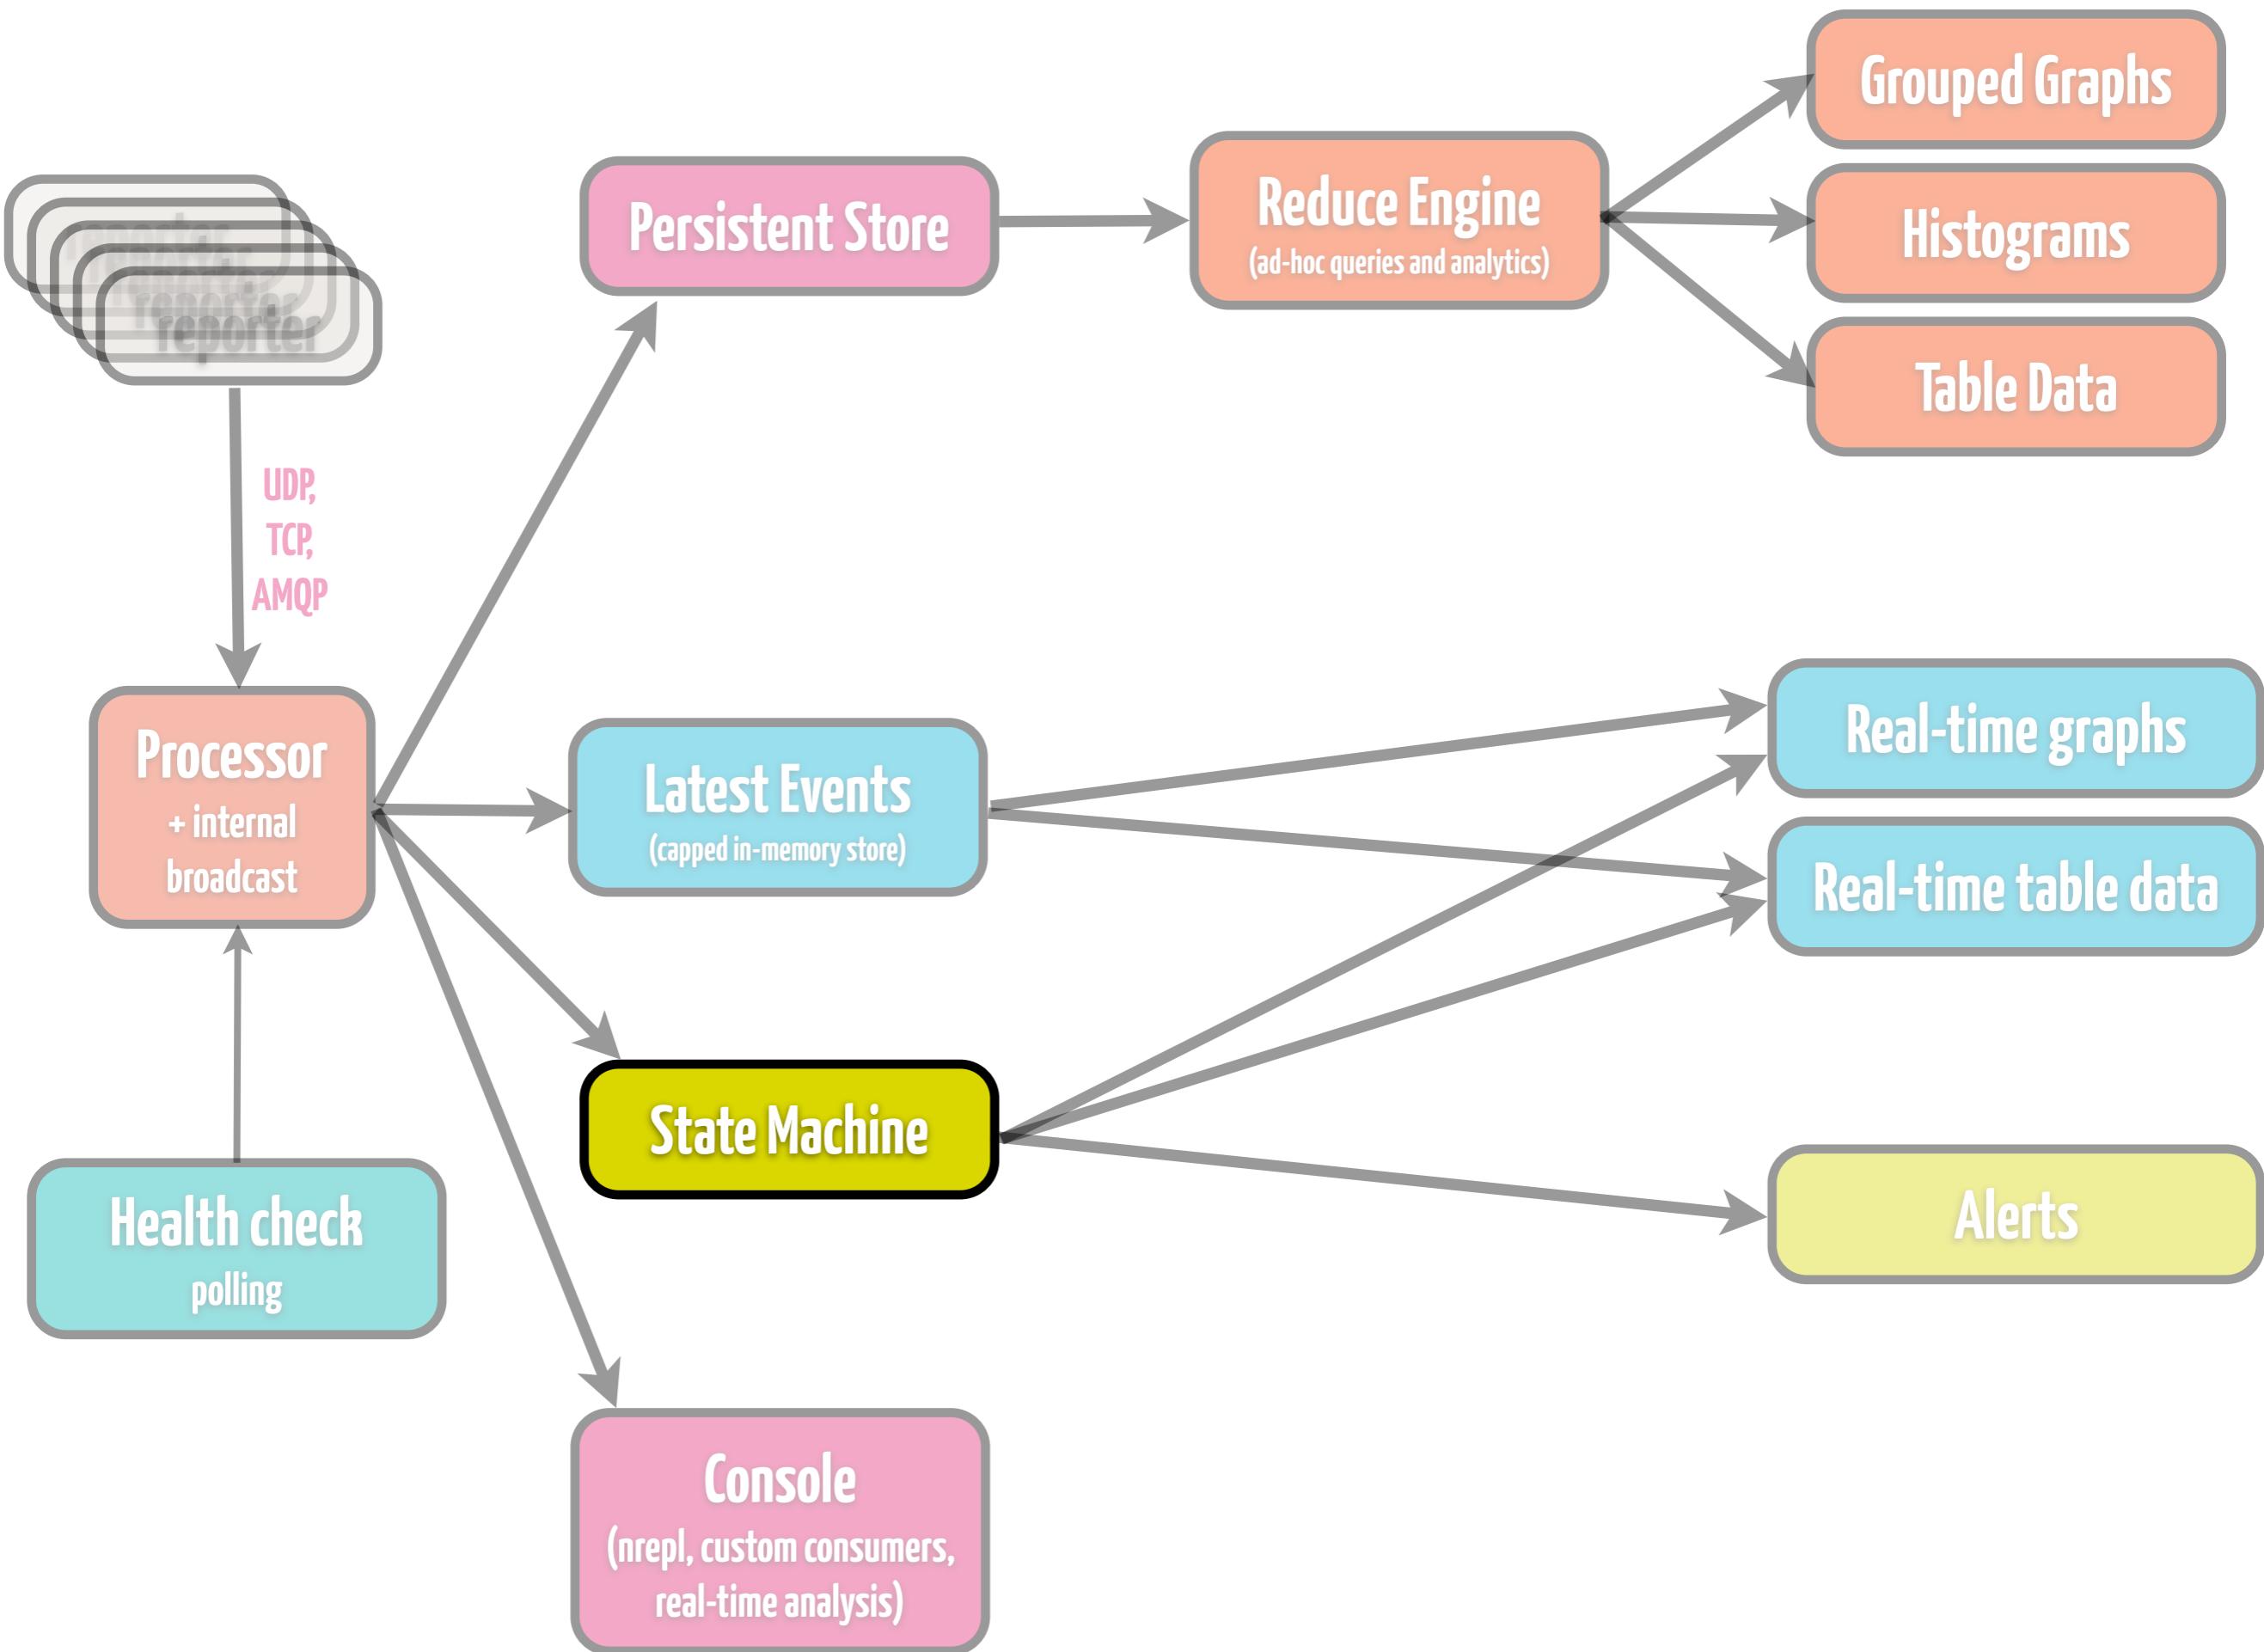
- that's basically an event cache.
- ring-buffer, whose size is limited by rules
- either time-bound (last hour, 6 hours, 24 hours)
- or size-bound (max 10K events)
- Riemann's indexes are Cliff Click NonBlockingHashMap backed
- you can go with (atom {:keyword (ref [])}) for starters
- fire up an nrepl server on collector
- have complete clojure access to your live data

Latest Events
(capped in-memory store)

- very queriable (100K+ events in collection are processed blazingly fast)
- optimize, calculate real-time indexes when needed
- for the most complex queries, that turn out to be a bottleneck, use state machine
- use for anything: stats, graphs, table data, dashboards. It's also very, very easy!

Latest Events
(capped in-memory store)

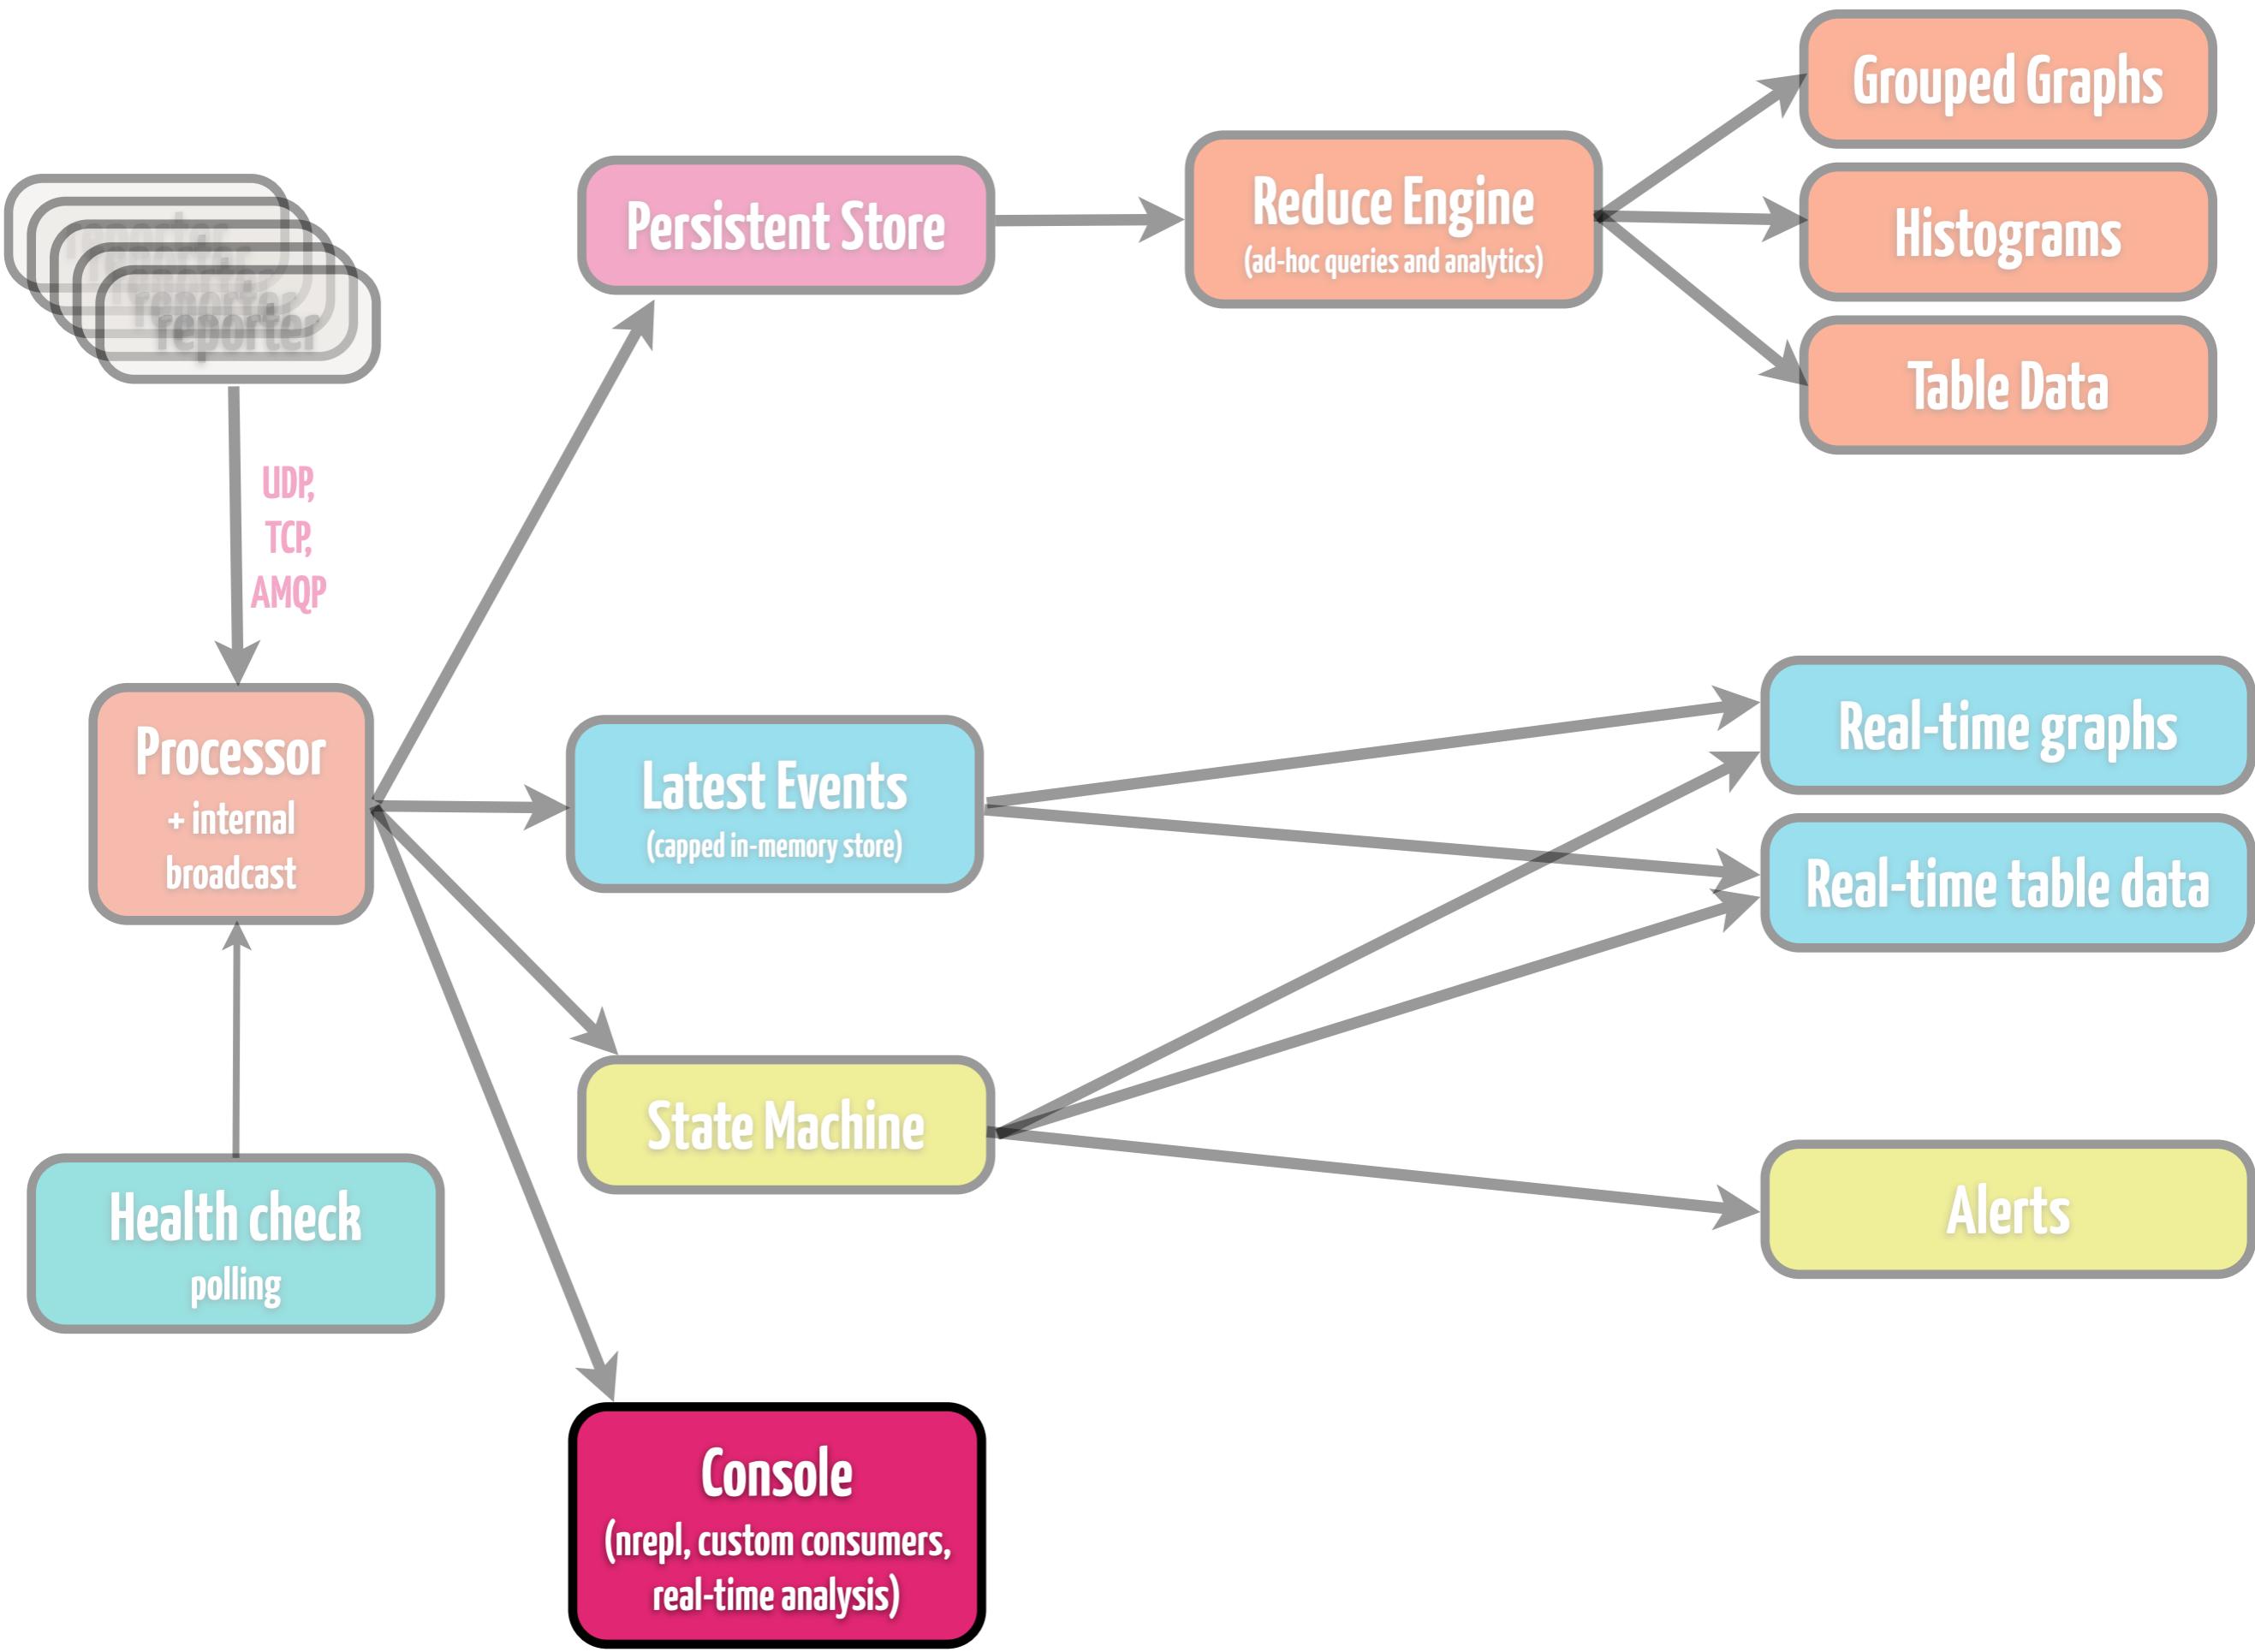
- very queriable (100K+ events in collection are processed blazingly fast)
- optimize, calculate real-time indexes when needed
- for the most complex queries, that turn out to be a bottleneck, use state machine
- use for anything: stats, graphs, table data, dashboards. It's also very, very easy!



- counters, gauges, meters
 - triggers for events and alarms
 - reducers of any type, actually
-
- if the amount of events of type `T` with md5 `M` (or any md5) exceeds 10, escalate the issue (send alert)
 - if we received more than `N` events of type `T` in 1 minute, send alert
 - if response of page exceeds `X` ms, send alert

State Machine

- You can measure metrics based on any arbitrary rule and make any arbitrary callback, basically.



Console
(nrepl, custom consumers,
real-time analysis)

- build processing pipelines in runtime
- run nrepl
- run queries, get necessary information

it's important to
know

what your
application does
right now

make it easy to
notice failures

to identify
behavior oddities

for that, again

you have to run a
lot of ad-hoc
queries

make system

more stable

by noticing

things that

happen

often

but irregular

see what code
doesn't reveal

if the host is up,

is “up” what you

expect by “up”?

collect
everything

storage is cheap

**YMMV if
if you operate on
twitter/facebook
scale, where you need to
seek more global approaches**

let your system
talk to you

extract business
metrics from
tech data

improve business
metrics by tech
data

make it easy to
see what your
users see

plan your growth

probably quite hard to do it without knowing your current corellation between performance and number of servers, and extracting trends from historical data

DYI

I've set it already quite a few times, but I won't get tired of saying it: DO IT YOURSELF.
it's possible to find a tool that does most of things for you, but you may not find your own patterns
with it. Use libs for visualization, persistency, distribution, transport, processing, but try to avoid
coupling. If it's hard to get data in, or get data out of it, good idea to avoid it.

also, do you know about

clojure



Wait for it...

Citytv

Werbz



It just werkz

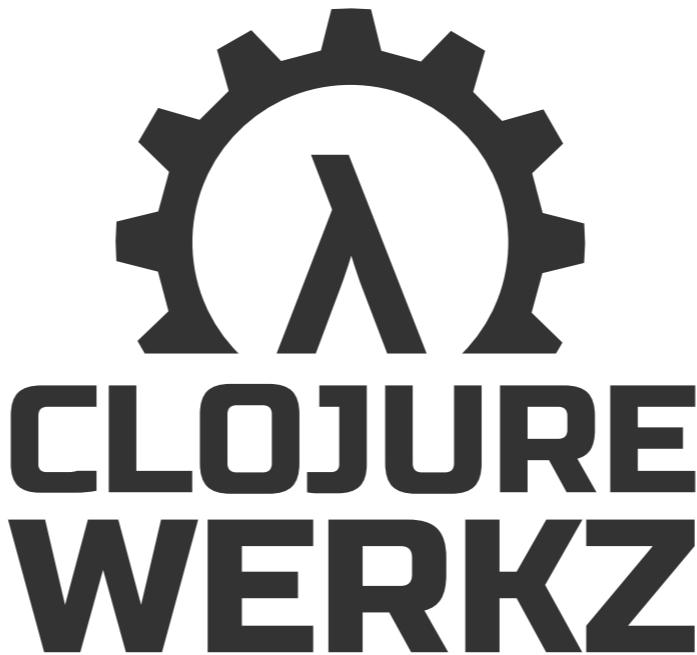
A growing collection of open source Clojure libraries that support multiple Clojure & JDK versions, licensed under the [EPL](#), target Clojure 1.3+.



How awesome is that? **Citytv**



That's All Folks!



 /  **ifesdjeen**

 /  **clojurewerkz**

Also, you should follow

 /  **michaelklishin**